

10-605/805 – ML for Large Datasets

Lecture 5: Distributed Linear Regression

Henry Chai

9/13/22

Front Matter

- HW1 released 8/30, due 9/13 (today!) at 11:59 PM
 - Recitation 3 on 9/16 will go over HW1 solutions
- HW2 released 9/8, due 9/22 at 11:59 PM
- Mini-project details released 9/9

Background: Big O Notation

- Used to describe an algorithm's time or space (storage) complexity in terms of the input size
- Formally:
$$f(x) = O(g(x)) \iff \exists C, x_0 \text{ s.t. } f(x) \leq Cg(x) \forall x \geq x_0$$
 - $O(1)$ = constant time/space, i.e., a fixed number of operations or storage regardless of input
 - $O(\log(n))$ = logarithmic time/space
 - $O(n)$ = linear time/space
- An algorithm's time and space complexity can be different
 - Example: multiplying an $a \times b$ matrix with an $b \times c$ matrix takes $O(abc)$ time (ac dot products between b -length vectors) but the result uses $O(ac)$ storage

Background: Empirical Risk Minimization

- A common framework for supervised learning
- Given:
 - some labelled training dataset $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^n$
 - a loss function $\ell: \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$
 - a hypothesis class or set of functions \mathcal{F}

the goal is to find

$$\hat{f} = \operatorname{argmin}_{f \in \mathcal{F}} \sum_{i=1}^n \ell(f(\mathbf{x}^{(i)}), y^{(i)})$$

with the hope that

$$\mathbb{E}_{p(\mathbf{x}, y)}[\ell(f(\mathbf{x}), y)] \approx \sum_{i=1}^n \ell(f(\mathbf{x}^{(i)}), y^{(i)})$$

Background: Empirical Risk Minimization

- A common framework for supervised learning
- Given:
 - some labelled training dataset $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^n$
 - a loss function $\ell: \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$
 - a hypothesis class or set of functions \mathcal{F}

the goal is to find

$$\hat{f} = \operatorname{argmin}_{f \in \mathcal{F}} \sum_{i=1}^n \ell(f(\mathbf{x}^{(i)}), y^{(i)})$$

- Depending on the choice of \mathcal{F} and ℓ , this objective function may be convex (easy to optimize) or non-convex (hard)
- Our focus will be solving this problem for large n and/or k

Background: Regression

- A type of supervised learning
- Given:
 - some labelled training dataset $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^n$
 - a loss function $\ell: \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$ where $\mathcal{Y} = \mathbb{R}$
 - a hypothesis class or set of functions \mathcal{F}

the goal is to find

$$\hat{f} = \operatorname{argmin}_{f \in \mathcal{F}} \sum_{i=1}^n \ell(f(\mathbf{x}^{(i)}), y^{(i)})$$

- Fun example: predicting the year a song was released based on (a representation of) its audio (see HW2)

Background: Linear Regression (Ordinary Least Squares)

- A type of supervised learning
 - Given:
 - some labelled training dataset $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^n$
 - $\ell(y, y') = (y - y')^2$
 - \mathcal{F} = all functions of the form $f(\mathbf{x}) = w_0 + \underbrace{\sum_{d=1}^k w_d x_d}$
- the goal is to find

$$\hat{f} = \operatorname{argmin}_{f \in \mathcal{F}} \sum_{i=1}^n (f(\mathbf{x}^{(i)}) - y^{(i)})^2$$

Background: Linear Regression (Ordinary Least Squares)

- A type of supervised learning

- Given:

- some labelled training dataset $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^n$
- $\ell(y, y') = (y - y')^2$
- \mathcal{F} = all functions of the form $f(\mathbf{x}) = \mathbf{w}^T \underbrace{[1 \ \mathbf{x}^T]^T}_{\{\omega_0, \omega_1, \dots, \omega_k\}}$

the goal is to find

$$\hat{f} = \operatorname{argmin}_{f \in \mathcal{F}} \sum_{i=1}^n (f(\mathbf{x}^{(i)}) - y^{(i)})^2$$

Background: Linear Regression (Ordinary Least Squares)

- A type of supervised learning
 - Given:
 - some labelled training dataset $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^n$
 - $\ell(y, y') = (y - y')^2$
 - \mathcal{F} = all functions of the form $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$
- the goal is to find
- \mathbf{w} implicitly prepended

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \sum_{i=1}^n (\mathbf{w}^T \mathbf{x}^{(i)} - y^{(i)})^2$$


Interpretable

Very easy to
optimize

- exact solution

Simple - tend to
generalize better

Fast inference



Background: Linear Regression (Ordinary Least Squares)

- A type of supervised learning
 - Given:
 - some labelled training dataset $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^n$
 - $\ell(y, y') = (y - y')^2$
 - \mathcal{F} = all functions of the form $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$
- the goal is to find
- 1 implicitly prepended

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \sum_{i=1}^n (\mathbf{w}^T \mathbf{x}^{(i)} - y^{(i)})^2$$

Background: Linear Regression (Ordinary Least Squares)

- A type of supervised learning
- Given:
 - some labelled training dataset $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^n$
 - $\ell(y, y') = (y - y')^2$
 - \mathcal{F} = all functions of the form $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$
the goal is to find 1 implicitly prepended

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{X}\mathbf{w} - \mathbf{y})^T (\mathbf{X}\mathbf{w} - \mathbf{y})$$

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \underbrace{\|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2}$$

- where $\mathbf{X} = \begin{bmatrix} \mathbf{x}^{(1)T} \\ \vdots \\ \mathbf{x}^{(n)T} \end{bmatrix}$ and $\mathbf{y} = \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(n)} \end{bmatrix}$

Background:
Linear
Regression
(Ordinary Least
Squares)

$$L_{\mathcal{D}}(\mathbf{w}) = (\mathbf{X}\mathbf{w} - \mathbf{y})^T (\mathbf{X}\mathbf{w} - \mathbf{y})$$

$$= \mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - 2\mathbf{y}^T \mathbf{X}^T \mathbf{w} + \mathbf{y}^T \mathbf{y}$$

$$\nabla_{\mathbf{w}} L_{\mathcal{D}}(\mathbf{w}) = 2\mathbf{X}^T \mathbf{X} \mathbf{w} - 2\mathbf{X}^T \mathbf{y} + \mathbf{0}$$

$$2\mathbf{X}^T \mathbf{X} \hat{\mathbf{w}} - 2\mathbf{X}^T \mathbf{y} = \mathbf{0}$$

$$\Rightarrow 2\mathbf{X}^T \mathbf{X} \hat{\mathbf{w}} = 2\mathbf{X}^T \mathbf{y}$$

$$\Rightarrow \hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

$$H_{\mathcal{D}} L_{\mathcal{D}}(\mathbf{w}) = 2\mathbf{X}^T \mathbf{X} \text{ which is positive semi-definite everywhere}$$

Background: Regularization

- A modification to empirical risk minimization that penalizes model complexity in order to combat overfitting
- Given:
 - some labelled training dataset $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^n$
 - a loss function $\ell: \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$
 - a hypothesis class or set of functions \mathcal{F}
 - a regularizer $R: \mathcal{W} \rightarrow \mathbb{R}$
 - a coefficient of regularization λ

the goal is to find

$$\hat{\mathbf{w}} = \operatorname{argmin}_{\mathbf{w}} \underbrace{\sum_{i=1}^n \ell(f_{\mathbf{w}}(\mathbf{x}^{(i)}), y^{(i)})}_{\text{empirical risk}} + \underbrace{\lambda R(\mathbf{w})}_{\text{regularizer}}$$

Background: Ridge Regression

- A modification to empirical risk minimization that penalizes model complexity in order to combat overfitting
- Given:
 - some labelled training dataset $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^n$
 - $\ell(y, y') = (y - y')^2$
 - \mathcal{F} = all functions of the form $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$
 - $R(\mathbf{w}) = \underbrace{\|\mathbf{w}\|_2^2}_{\mathbf{w}^T \mathbf{w}}$
 - a coefficient of regularization λ

the goal is to find

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{X}\mathbf{w} - \mathbf{y})^T (\mathbf{X}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w}$$

Background: Ridge Regression

$$L_{\mathcal{D}}(\mathbf{w}) = (\mathbf{X}\mathbf{w} - \mathbf{y})^T (\mathbf{X}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w}$$

•
•
•

$$\rightarrow \hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}_k)^{-1} \mathbf{X}^T \mathbf{y}$$

where \mathbf{I}_k is the $k \times k$ identity matrix

Aside:
How can we
set λ ?

$$L_{\mathcal{D}}(\mathbf{w}) = (\mathbf{X}\mathbf{w} - \mathbf{y})^T (\mathbf{X}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w}$$

•
•
•

$$\rightarrow \hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}_k)^{-1} \mathbf{X}^T \mathbf{y}$$

where \mathbf{I}_k is the $k \times k$ identity matrix

$$(\lambda > 0)$$

Recall: Machine Learning Pipeline

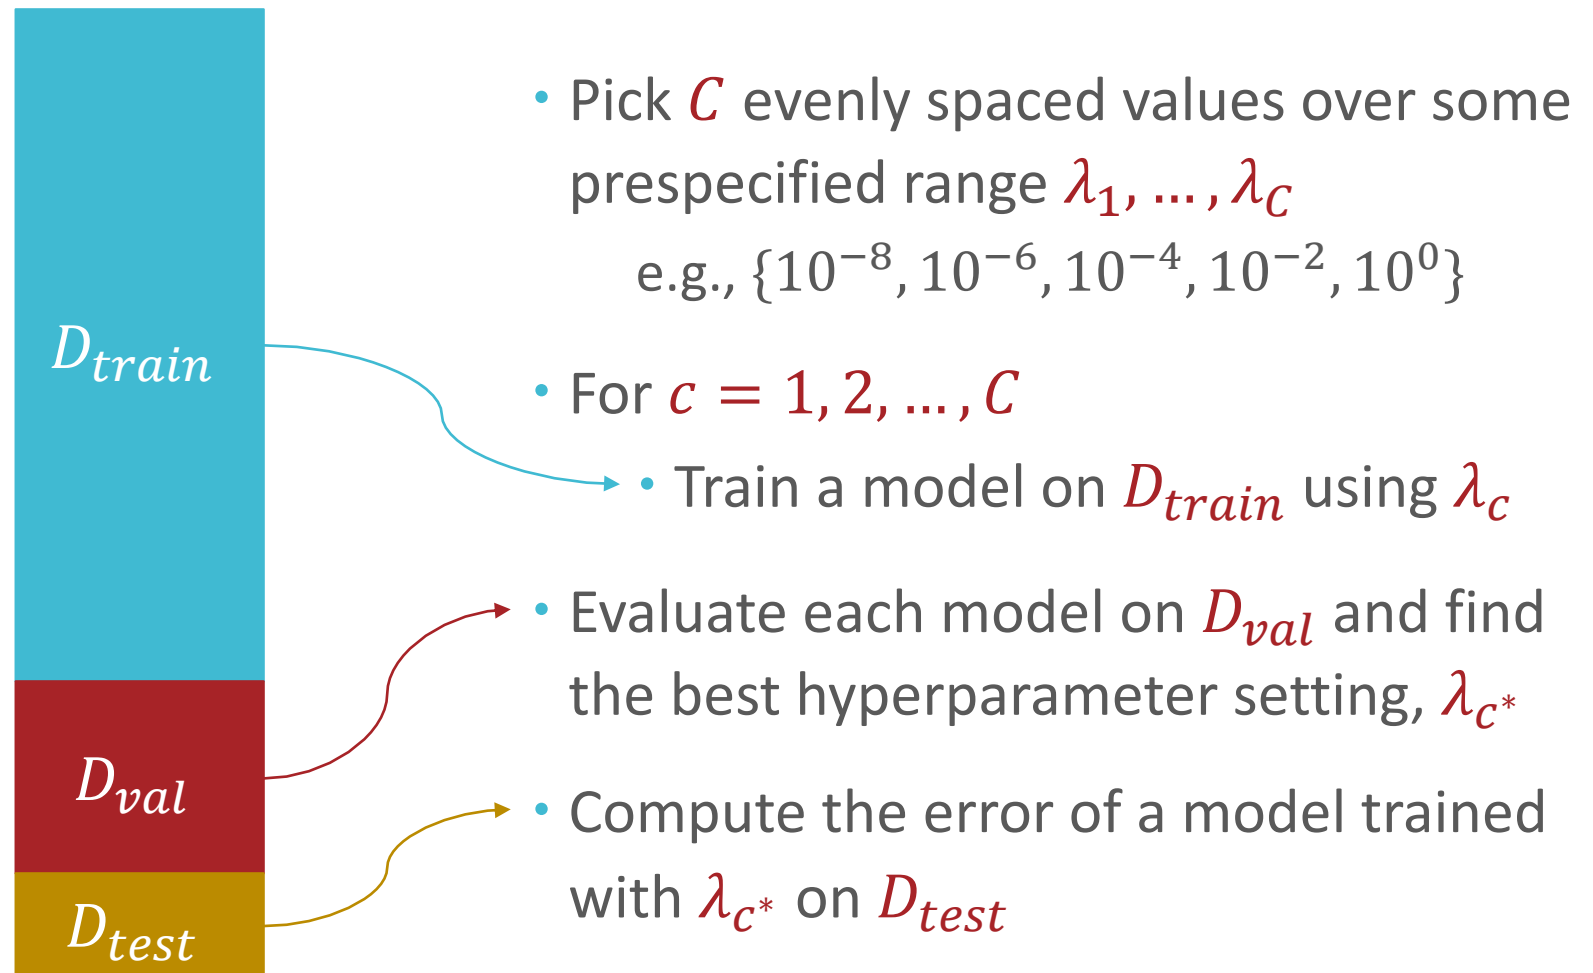
- Hyperparameter optimization



- Suppose we want to compare multiple hyperparameter settings $\theta_1, \dots, \theta_C$ (where do these come from???)
- For $c = 1, 2, \dots, C$
 - Train a model on D_{train} using θ_c
 - Evaluate each model on D_{val} and find the best hyperparameter setting, θ_{c^*}
 - Compute the error of a model trained with θ_{c^*} on D_{test}

HW2 Preview: Grid Search

- Hyperparameter optimization



Linear Regression: Computational Cost

$$\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

1. Does this quantity exist, i.e., is $\mathbf{X}^T \mathbf{X}$ invertible?

For the most part yes, $\mathbf{X}^T \mathbf{X}$ is positive semi-definite; regularization helps

2. If so, how expensive is it to compute?

$$\mathbf{X} \in \mathbb{R}^{n \times k} \Rightarrow \mathbf{X}^T \mathbf{X} \in \mathbb{R}^{k \times k}$$

Time cost of inverting $\mathbf{X}^T \mathbf{X} = O(k^3)$

Space cost of inverting $\mathbf{X}^T \mathbf{X} = O(k^2)$

Linear Regression: Large n , Small k

- Assume $O(k^3)$ computation and $O(k^2)$ storage is possible on a single machine

✓ We can store and invert $X^T X$

We cannot compute $X^T X$ } on a single machine

We cannot store X

- Idea: distribute storage of X and computation of $X^T X$
 1. Store the rows of X across different machines
 2. Compute $X^T X$ as the sum of outer products

Matrix Multiplication via Inner Products

$$\begin{bmatrix} 1 & 4 & 5 \\ 3 & 1 & 3 \end{bmatrix} \begin{bmatrix} 2 \\ 2 \\ 3 \end{bmatrix} = \begin{bmatrix} \\ \end{bmatrix}$$

$$1 * 2 + 4 * 2 + 5 * 3 = 25$$

Matrix Multiplication via Inner Products


$$\begin{bmatrix} 1 & 4 & 5 \\ 3 & 1 & 3 \end{bmatrix} \begin{bmatrix} 2 \\ 2 \\ 4 \end{bmatrix} = \begin{bmatrix} 25 \\ \end{bmatrix}$$

$$1 * 2 + 4 * 1 + 5 * 4 = 26$$

Matrix Multiplication via Inner Products

$$\begin{bmatrix} 1 & 4 & 5 \\ 3 & 1 & 3 \end{bmatrix} \begin{bmatrix} 2 & 2 \\ 2 & 1 \\ 3 & 4 \end{bmatrix} = \begin{bmatrix} 25 & 26 \\ 17 & 19 \end{bmatrix}$$

Matrix Multiplication via Inner Products



Matrix Multiplication via Outer Products

$$\begin{bmatrix} 1 & 4 & 5 \\ 3 & 1 & 3 \end{bmatrix} \begin{bmatrix} 2 & 2 \\ 1 & 4 \end{bmatrix} = \begin{bmatrix} 2 & 2 \\ 6 & 6 \end{bmatrix}$$

Matrix Multiplication via Outer Products

$$\begin{bmatrix} 1 & 4 & 5 \\ 3 & 1 & 3 \end{bmatrix} \begin{bmatrix} 2 & 2 \\ 2 & 1 \\ 3 & 4 \end{bmatrix} = \begin{bmatrix} & & \\ & & \\ & & \end{bmatrix}$$

$$\begin{bmatrix} 2 & 2 \\ 6 & 6 \end{bmatrix} + \begin{bmatrix} 8 & 4 \\ 2 & 1 \end{bmatrix}$$

Matrix Multiplication via Outer Products

$$\begin{bmatrix} 1 & 4 & 5 \\ 3 & 1 & 3 \end{bmatrix} \begin{bmatrix} 2 & 2 \\ 2 & 1 \\ 3 & 4 \end{bmatrix} = \begin{bmatrix} & \\ & \end{bmatrix}$$

$$\begin{bmatrix} 2 & 2 \\ 6 & 6 \end{bmatrix} + \begin{bmatrix} 8 & 4 \\ 2 & 1 \end{bmatrix} + \begin{bmatrix} 15 & 20 \\ 9 & 12 \end{bmatrix}$$

Matrix Multiplication via Outer Products

$$\begin{bmatrix} 1 & 4 & 5 \\ 3 & 1 & 3 \end{bmatrix} \begin{bmatrix} 2 & 2 \\ 2 & 1 \\ 3 & 4 \end{bmatrix} = \begin{bmatrix} 25 & 26 \\ 17 & 19 \end{bmatrix}$$

$$\begin{bmatrix} 2 & 2 \\ 6 & 6 \end{bmatrix} + \begin{bmatrix} 8 & 4 \\ 2 & 1 \end{bmatrix} + \begin{bmatrix} 15 & 20 \\ 9 & 12 \end{bmatrix}$$

Matrix Multiplication via Outer Products

$$\begin{bmatrix} A_{11} & A_{12} & \cdots & A_{1m} \\ A_{21} & A_{22} & \cdots & A_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ A_{n1} & A_{n2} & \cdots & A_{nm} \end{bmatrix} \begin{bmatrix} B_{11} & \cdots & B_{1k} \\ B_{21} & \cdots & B_{2k} \\ B_{31} & \cdots & B_{3k} \\ \vdots & \ddots & \vdots \\ B_{m1} & \cdots & B_{mk} \end{bmatrix}$$

$$= \begin{bmatrix} \sum_{i=1}^m A_{1i} B_{i1} & \cdots & \sum_{i=1}^m A_{1i} B_{ik} \\ \vdots & \ddots & \vdots \\ \sum_{i=1}^m A_{ni} B_{i1} & \cdots & \sum_{i=1}^m A_{ni} B_{ik} \end{bmatrix} = \sum_{i=1}^m \begin{bmatrix} A_{1i} B_{i1} & \cdots & A_{1i} B_{ik} \\ \vdots & \ddots & \vdots \\ A_{ni} B_{i1} & \cdots & A_{ni} B_{ik} \end{bmatrix}$$

Distributed Computation of $(X^T X)^{-1}$

$$X^T X = \begin{bmatrix} \uparrow & \uparrow & \dots & \uparrow \\ \mathbf{x}^{(1)} & \mathbf{x}^{(2)} & \dots & \mathbf{x}^{(n)} \\ \downarrow & \downarrow & \dots & \downarrow \end{bmatrix} \begin{bmatrix} \leftarrow & \mathbf{x}^{(1)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(2)T} & \rightarrow \\ \vdots & \vdots & \vdots \\ \leftarrow & \mathbf{x}^{(n)T} & \rightarrow \end{bmatrix} = \sum_{i=1}^n \mathbf{x}^{(i)} \mathbf{x}^{(i)T}$$

- Idea: distribute $\mathbf{x}^{(i)}$ and compute summands in parallel

Workers

$$\begin{bmatrix} \leftarrow & \mathbf{x}^{(1)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(4)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$$

$$\begin{bmatrix} \leftarrow & \mathbf{x}^{(2)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(3)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$$

$$\begin{bmatrix} \leftarrow & \mathbf{x}^{(5)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(7)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$$

Distributed Computation of $(X^T X)^{-1}$

Workers	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(1)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(4)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(2)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(3)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(5)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(7)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$
Map	$\mathbf{x}^{(i)} \mathbf{x}^{(i)T}$	$\mathbf{x}^{(i)} \mathbf{x}^{(i)T}$	$\mathbf{x}^{(i)} \mathbf{x}^{(i)T}$

Distributed Computation of $(X^T X)^{-1}$

Workers	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(1)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(4)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(2)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(3)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(5)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(7)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$
Map	$\mathbf{x}^{(i)} \mathbf{x}^{(i)T}$	$\mathbf{x}^{(i)} \mathbf{x}^{(i)T}$	$\mathbf{x}^{(i)} \mathbf{x}^{(i)T}$
Reduce	$\left(\sum_{i=1}^n \mathbf{x}^{(i)} \mathbf{x}^{(i)T} \right)^{-1}$		

Distributed Computation of $(X^T X)^{-1}$

Workers	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(1)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(4)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(2)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(3)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(5)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(7)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$
Map	$\mathbf{x}^{(i)} \mathbf{x}^{(i)T}$	$\mathbf{x}^{(i)} \mathbf{x}^{(i)T}$	$\mathbf{x}^{(i)} \mathbf{x}^{(i)T}$
Reduce	$\left(\sum_{i=1}^n \mathbf{x}^{(i)} \mathbf{x}^{(i)T} \right)^{-1}$		

$O(nk)$ distributed
storage (total)

Distributed Computation of $(X^T X)^{-1}$

Workers	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(1)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(4)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(2)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(3)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(5)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(7)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$O(nk)$ distributed storage (total)	
Map	$\mathbf{x}^{(i)} \mathbf{x}^{(i)T}$	$\mathbf{x}^{(i)} \mathbf{x}^{(i)T}$	$\mathbf{x}^{(i)} \mathbf{x}^{(i)T}$	$O(nk^2)$ distributed work (total)	$O(k^2)$ local storage
Reduce	$\left(\sum_{i=1}^n \mathbf{x}^{(i)} \mathbf{x}^{(i)T} \right)^{-1}$				

Distributed Computation of $(X^T X)^{-1}$

Workers	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(1)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(4)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(2)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(3)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(5)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(7)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$O(nk)$ distributed storage (total)	
Map	$\mathbf{x}^{(i)} \mathbf{x}^{(i)T}$	$\mathbf{x}^{(i)} \mathbf{x}^{(i)T}$	$\mathbf{x}^{(i)} \mathbf{x}^{(i)T}$	$O(nk^2)$ distributed work (total)	$O(k^2)$ local storage
Reduce	$\left(\sum_{i=1}^n \mathbf{x}^{(i)} \mathbf{x}^{(i)T} \right)^{-1}$			$O(k^3)$ local work	$O(k^2)$ local storage

Distributed Computation of $(X^T X)^{-1}$

Workers	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(1)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(4)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(2)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(3)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(5)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(7)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$O(nk)$ distributed storage (total)	
<code>trainData.map(compute_outer_prods)</code>				$O(nk^2)$ distributed work (total)	$O(k^2)$ local storage
<code>trainData.reduce(sum_and_invert)</code>				$O(k^3)$ local work	$O(k^2)$ local storage

Distributed Computation of $(X^T X)^{-1}$

Linear Regression: Large n , Large k

- Now, $O(k^3)$ computation and $O(k^2)$ storage is *not* possible on a single machine

We cannot store and invert $X^T X$

We cannot compute $X^T X$


We cannot store X

10-605/805 Principle #1: computation and storage should be at most linear in n and k

- Idea: use a different algorithm!


Background: Gradient Descent

- An iterative method for minimizing functions
- Requires the gradient to exist everywhere




Background: Gradient Descent

- An iterative method for minimizing functions
- Requires the gradient to exist everywhere




Background: Gradient Descent

- An iterative method for minimizing functions
- Requires the gradient to exist everywhere



Background: Gradient Descent

- An iterative method for minimizing functions
- Requires the gradient to exist everywhere



- Good news: the linear regression objective is convex so gradient descent will always converge to the global minimum


Background: Gradient Descent

- Suppose we're trying to minimize some function L and we're currently at some location $\mathbf{w}^{(t)}$
- Move some distance, α , in the “most downhill” direction, \mathbf{v} :
$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} + \alpha \mathbf{v}$$
- The gradient points in the direction of steepest *increase* ...
- ... so let's move in the opposite direction!

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \alpha \nabla_{\mathbf{w}} L(\mathbf{w}^{(t)})$$


Background: Gradient Descent

Direction of
gradient



$$2x < 0 \text{ for } x < 0$$

Direction of
gradient




$$2x > 0 \text{ for } x > 0$$

$$\frac{\partial}{\partial x} x^2 = 2x$$


Background: Gradient Descent

Direction of
global minimum



$$2x < 0 \text{ for } x < 0$$


Direction of
global minimum




$$2x > 0 \text{ for } x > 0$$

$$\frac{\partial}{\partial x} x^2 = 2x$$

Background: Gradient Descent Step Size




Small α




Large α

Background: Gradient Descent Step Size




Small α




Large α

Background: Gradient Descent Step Size




Small α



Large α

Background: Gradient Descent Step Size

- Use a variable $\alpha^{(t)}$ instead of a fixed α !



- Example: $\alpha^{(t)} = \frac{\alpha}{n\sqrt{t}}$

Gradient Descent for Linear Regression

- Input: $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^n, \alpha$
- 1. Initialize $\mathbf{w}^{(0)}$ to all zeros and set $t = 0$
- 2. While TERMINATION CRITERION is not satisfied
 - a. Compute the gradient:

$$\nabla_{\mathbf{w}} L_{\mathcal{D}}(\mathbf{w}^{(t)}) = \underline{(2X^T X \mathbf{w}^{(t)} - 2X^T \mathbf{y})}$$

- b. Update the weights:

$$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \underset{\uparrow}{\frac{\alpha}{n\sqrt{t}}} (2X^T X \mathbf{w}^{(t)} - 2X^T \mathbf{y})$$

- c. Increment t : $t \leftarrow t + 1$

- Output: $\mathbf{w}^{(t)}$

Gradient Descent for Linear Regression

- Input: $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^n, \alpha$

1. Initialize $\mathbf{w}^{(0)}$ to all zeros and set $t = 0$
2. While TERMINATION CRITERION is not satisfied
 - a. Compute the gradient:

$$\nabla_{\mathbf{w}} L_{\mathcal{D}}(\mathbf{w}^{(t)}) = 2 \sum_{i=1}^n (\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$$

- b. Update the weights:

$$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \frac{\alpha}{n\sqrt{t}} \sum_{i=1}^n (\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$$

- c. Increment t : $t \leftarrow t + 1$

- Output: $\mathbf{w}^{(t)}$

Idea: distribute $\mathbf{x}^{(i)}$ and compute summands in parallel

- Input: $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^n, \alpha$

1. Initialize $\mathbf{w}^{(0)}$ to all zeros and set $t = 0$
2. While TERMINATION CRITERION is not satisfied

- a. Compute the gradient:

$$\nabla_{\mathbf{w}} L_{\mathcal{D}}(\mathbf{w}^{(t)}) = 2 \underbrace{\sum_{i=1}^n (\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}}_{\text{gradient}}$$

- b. Update the weights:

$$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \frac{\alpha}{n\sqrt{t}} \sum_{i=1}^n (\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$$

- c. Increment t : $t \leftarrow t + 1$

- Output: $\mathbf{w}^{(t)}$

Workers	$\begin{bmatrix} \leftarrow \mathbf{x}^{(1)T} \rightarrow \\ \leftarrow \mathbf{x}^{(4)T} \rightarrow \\ \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow \mathbf{x}^{(2)T} \rightarrow \\ \leftarrow \mathbf{x}^{(3)T} \rightarrow \\ \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow \mathbf{x}^{(5)T} \rightarrow \\ \leftarrow \mathbf{x}^{(7)T} \rightarrow \\ \vdots \end{bmatrix}$
Map	$(\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$	$(\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$	$(\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$
Reduce	$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \frac{\alpha}{n\sqrt{t}} \sum_{i=1}^n (\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$		

Distributed Gradient Descent

Workers	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(1)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(4)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(2)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(3)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(5)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(7)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$
Map	$(\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$	$(\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$	$(\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$
Reduce	$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \frac{\alpha}{n\sqrt{t}} \sum_{i=1}^n (\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$		

$O(nk)$ distributed storage (total)

Distributed Gradient Descent

Workers	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(1)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(4)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(2)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(3)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(5)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(7)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$O(nk)$ distributed storage (total)	
Map	$(\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$	$(\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$	$(\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$	$O(nk)$ distributed work (total)	$O(k)$ local storage
Reduce	$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \frac{\alpha}{n\sqrt{t}} \sum_{i=1}^n (\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$				

Distributed Gradient Descent

Workers	$\begin{bmatrix} \leftarrow \mathbf{x}^{(1)T} \rightarrow \\ \leftarrow \mathbf{x}^{(4)T} \rightarrow \\ \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow \mathbf{x}^{(2)T} \rightarrow \\ \leftarrow \mathbf{x}^{(3)T} \rightarrow \\ \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow \mathbf{x}^{(5)T} \rightarrow \\ \leftarrow \mathbf{x}^{(7)T} \rightarrow \\ \vdots \end{bmatrix}$	$O(nk)$ distributed storage (total)	
Map	$(\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$	$(\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$	$(\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$	$O(nk)$ distributed work (total)	$O(k)$ local storage
Reduce	$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \frac{\alpha}{n\sqrt{t}} \sum_{i=1}^n (\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$			$O(k)$ local work	$O(k)$ local storage

Distributed Gradient Descent

Workers	$\begin{bmatrix} \leftarrow \mathbf{x}^{(1)T} \rightarrow \\ \leftarrow \mathbf{x}^{(4)T} \rightarrow \\ \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow \mathbf{x}^{(2)T} \rightarrow \\ \leftarrow \mathbf{x}^{(3)T} \rightarrow \\ \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow \mathbf{x}^{(5)T} \rightarrow \\ \leftarrow \mathbf{x}^{(7)T} \rightarrow \\ \vdots \end{bmatrix}$	$O(nk)$ distributed storage (total)	
Map	$(\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$	$(\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$	$(\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$	$O(nk)$ distributed work (total)	$O(k)$ local storage
Reduce	$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \frac{\alpha}{n\sqrt{t}} \sum_{i=1}^n (\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$			$O(k)$ local work	$O(k)$ local storage

Issue: all workers must have the latest weight vector

Distributed Gradient Descent

Workers

$$\begin{bmatrix} \leftarrow & \mathbf{x}^{(1)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(4)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$$

$$\begin{bmatrix} \leftarrow & \mathbf{x}^{(2)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(3)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$$

$$\begin{bmatrix} \leftarrow & \mathbf{x}^{(5)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(7)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$$

$O(nk)$ distributed storage (total)

`trainData.map(compute_pointwise_grads($\mathbf{w}^{(t)}$))`

$O(nk)$
distributed
work (total)

$O(k)$
local storage

`trainData.sum()`

$O(k)$
local work

$O(k)$
local storage

Issue: all workers must have the latest weight vector

Distributed Gradient Descent


Gradient Descent

- Pros:
 - Trivially parallelizable
 - Each individual iteration is cheap
 - Can be further improved using stochastic variants
 - Guaranteed to converge on convex objective functions
- Cons:
 - Potentially slow convergence
 - Introduction of a hyperparameter
 - Network communication in each iteration



Gradient Descent

- Pros:
 - Trivially parallelizable
 - Each individual iteration is cheap
 - Can be further improved using stochastic variants
 - Guaranteed to converge on convex objective functions
- Cons:
 - Potentially slow convergence
 - Introduction of a hyperparameter
 - **Network communication in each iteration**

Recall: Communication Hierarchy



10-605/805
Principle #2:
Perform parallel
and in-memory
computation
whenever
possible



10-605/805

Principle #2: Perform parallel and in-memory computation whenever possible

- Persisting data in-memory reduces communication, especially for iterative procedures

```
trainData.cache()    or    .persist()  
for t in range(num_iters):  
    alpha_t = alpha / n * sqrt(t)  
  
    grad = trainData.map(compute_pointwise_grads(w)).sum()  
  
    w -= alpha_t * grad
```


10-605/805

Principle #3: Minimize network communication

- Inherently at odds with Principle #2 → need to tradeoff between parallelism and network communication
- Three types of objects that may need to be communicated:
 - Data
 - Models
 - Intermediate objects
- Strategies:
 - Keep large objects local
 - Reduce the number of iterations

Workers	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(1)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(4)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(2)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(3)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & \mathbf{x}^{(5)T} & \rightarrow \\ \leftarrow & \mathbf{x}^{(7)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$O(nk)$ distributed storage (total)	
Map	$\mathbf{x}^{(i)} \mathbf{x}^{(i)T}$	$\mathbf{x}^{(i)} \mathbf{x}^{(i)T}$	$\mathbf{x}^{(i)} \mathbf{x}^{(i)T}$	$O(nk^2)$ distributed work (total)	$O(k^2)$ local storage
Reduce	$\left(\sum_{i=1}^n \mathbf{x}^{(i)} \mathbf{x}^{(i)T} \right)^{-1}$			$O(k^3)$ local work	$O(k^2)$ local storage

Data Parallel: Compute outer products locally

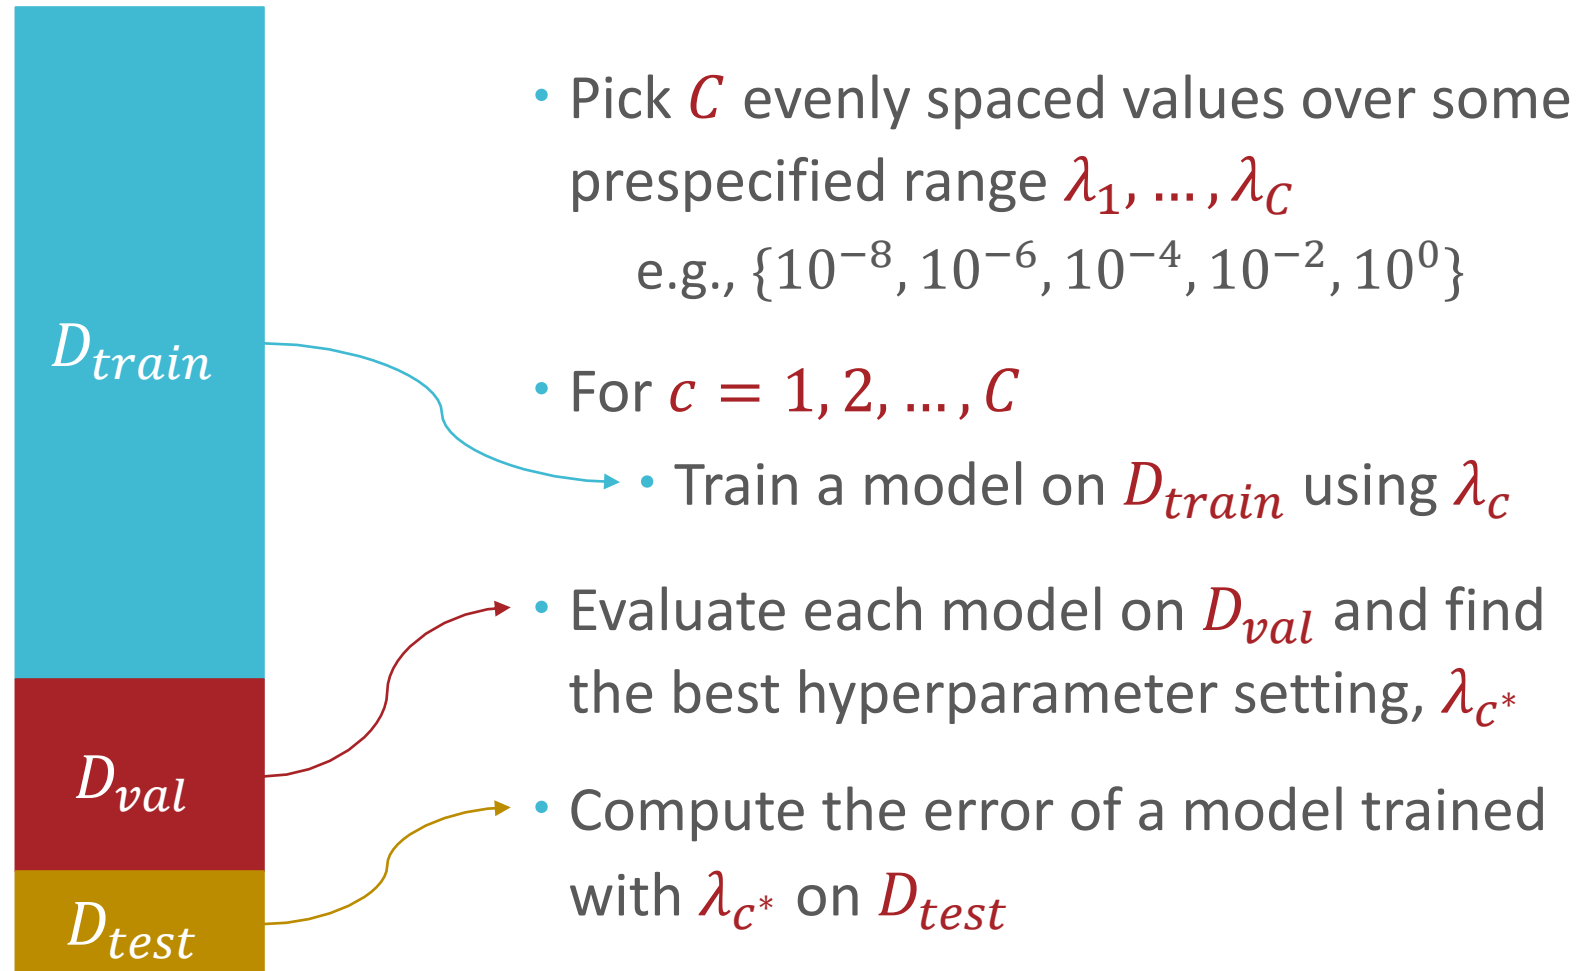
Workers	$\begin{bmatrix} \leftarrow \mathbf{x}^{(1)T} \rightarrow \\ \leftarrow \mathbf{x}^{(4)T} \rightarrow \\ \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow \mathbf{x}^{(2)T} \rightarrow \\ \leftarrow \mathbf{x}^{(3)T} \rightarrow \\ \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow \mathbf{x}^{(5)T} \rightarrow \\ \leftarrow \mathbf{x}^{(7)T} \rightarrow \\ \vdots \end{bmatrix}$	$O(nk)$ distributed storage (total)	
Map	$(\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$	$(\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$	$(\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$	$O(nk)$ distributed work (total)	$O(k)$ local storage
Reduce	$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \frac{\alpha}{n\sqrt{t}} \sum_{i=1}^n (\mathbf{w}^{(t)T} \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$			$O(k)$ local work	$O(k)$ local storage

Issue: all workers must have the latest weight vector

Data Parallel: Compute pointwise gradients locally

Model Parallel:
Train each
hyperparameter
setting on
different
machine(s)

- Hyperparameter optimization



Key Takeaways

- 10-605/805 Principles:
 1. Computation and storage should be linear in n and k
 - For linear regression:
 - When k is small, distribute covariance matrix computation using outer products
 - When k is large, minimize squared error via distributed gradient descent
 2. Perform parallel and in-memory computation whenever possible
 3. Minimize network communication
 - Data vs model parallelism