

# 10-605/805 – ML for Large Datasets

## Lecture 10: Randomized Algorithms

Henry Chai

9/29/22

# Front Matter

- HW3 released 9/23, due 10/4 at 11:59 PM
- Midterm exam on 10/11, two weeks from today!
  - Lecture on 10/6 (next Thursday) will be an *optional* practice exam - these will not be collected/graded
    - Recitation on 10/7 (next Friday) will go through the solutions
    - Both the practice exam and the solutions will be released after the Recitation
- Lecture on 10/4 (next Tuesday) will be an AWS tutorial; please bring your laptops to class on that day

# Two Specific Applications of Feature Hashing

1. *Count-min sketch* for logistic regression
2. *Locality sensitive hashing* for nearest neighbors/clustering

# Two Specific Applications of Feature Hashing

1. *Count-min sketch* for logistic regression
2. *Locality sensitive hashing* for nearest neighbors/clustering

# Count-min Sketch

- Data structure used to estimate the frequency of items in some stream of inputs
  - Running example: finding “viral” terms in Google searches (<https://trends.google.com/trends/?geo=US>)
  - Naïve approach: just keep an array with an index for every possible search term and add one to that index when the term appears
    - Massive array
    - Could be dynamic/need to grow if unseen search terms arrive

# Count-min Sketch: Connection to Logistic Regression

- If  $\mathbf{x}$  is a (sparse) one-hot encoded vector, then we can efficiently compute  $\mathbf{w}^{(t)T} \mathbf{x}$  as

$$\sum_{j: x_j \neq 0} w_j^{(t)} x_j \approx \sum_{j: x_j \neq 0} \left( \min_i C[i, h_i(j)] \right) x_j$$

1	2	3	...	$m$
$C[1,1]$	$C[1,2]$	$C[1,3]$	...	$C[1,m]$
$\vdots$	$\vdots$	$\vdots$	$\ddots$	$\vdots$
$C[r,1]$	$C[r,2]$	$C[r,3]$	...	$C[r,m]$

- This matrix holds a compact representation of  $\mathbf{w}^{(t)}$  at each iteration

# Count-min Sketch: Algorithm

- Input:  $r$  independent hash functions  $h_1, \dots, h_r$  that each map to  $m$  buckets
- 1. Initialize an  $r \times m$  count matrix,  $C$ , to all zeros
- 2. For each item,  $s$ , in some stream of data:
  - a. For  $i \in \{1, \dots, r\}$ :
    - i.  $C[i, h_i(s)] += 1$
- 3. For any item  $s$ , return  $\hat{c}_s = \min_i(C[i, h_i(s)])$  as an approximation  $c_s$ , the true number of occurrences of  $s$

Observation:  
Count-min  
sketch can only  
overestimate  
counts!

But by how  
much?

- Input:  $r$  independent hash functions  $h_1, \dots, h_r$  that each map to  $m$  buckets
  1. Initialize an  $r \times m$  count matrix,  $C$ , to all zeros
  2. For each item,  $s$ , in some stream of data:
    - a. For  $i \in \{1, \dots, r\}$ :
      - i.  $C[i, h_i(s)] += 1$
  3. For any item  $s$ , return  $\hat{c}_s = \min_i(C[i, h_i(s)])$  as an approximation  $c_s$ , the true number of occurrences of  $s$



# Count-min Sketch: Theoretical Guarantees

- Given  $r > \log_2 1/\delta$  hash functions that each map to  $m > 2/\epsilon$  buckets, then  $\forall s$

$$P(\hat{c}_s \geq c_s + \epsilon \|c\|_1) \leq \delta$$

where  $\|c\|_1$  is the total number of items in the stream.

- Key assumptions:
  - All hash functions are independent of each other
- $P(h_i(s) \mid h_j(s')) = P(h_i(s)) \forall i, j \in \{1, \dots, r\}$  and  $s, s'$
- Each hash function is *pairwise independent*

$$P(h_i(s) = b \cap h_i(s') = b') = \frac{1}{m^2} \forall i \in \{1, \dots, r\}, \\ b, b' \in \{0, \dots, m\} \text{ and } s, s'$$

$$\Rightarrow P(h_i(s) = h_i(s')) = \frac{1}{m} \forall i \in \{1, \dots, r\} \text{ and } s, s'$$

# Count-min Sketch: Theoretical Guarantees Proof

- Given  $r > \log_2 1/\delta$  hash functions that each map to  $m > 2/\epsilon$  buckets, then  $\forall s$

$$P(\hat{c}_s \geq c_s + \epsilon \|\mathbf{c}\|_1) \leq \delta$$

where  $\|\mathbf{c}\|_1$  is the total number of items in the stream.

- Proof:
  - For an arbitrary hash function,  $h_i$ , and item,  $s$ :

$$\mathbb{E}[C[i, h_i(s)]] = c_s + \sum_{s' \neq s} P(h_i(s) = h_i(s')) c_{s'}$$

(by pairwise independence)

$$= c_s + \sum_{s' \neq s} \frac{1}{m} c_{s'} = c_s + \frac{1}{m} \sum_{s' \neq s} c_{s'}$$

(by definition of  $m$  and  $\|\mathbf{c}\|_1$ )

$$\leq c_s + \frac{\epsilon}{2} \sum_{s' \neq s} c_{s'} \leq c_s + \frac{\epsilon}{2} \|\mathbf{c}\|_1$$

## Count-min Sketch: Theoretical Guarantees Proof (Cont.)

- Given  $r > \log_2 1/\delta$  hash functions that each map to  $m > 2/\epsilon$  buckets, then  $\forall s$

$$P(\hat{c}_s \geq c_s + \epsilon \|\mathbf{c}\|_1) \leq \delta$$

where  $\|\mathbf{c}\|_1$  is the total number of items in the stream.

- Proof:
    - For an arbitrary hash function,  $h_i$ , and item,  $s$ :
- $$P(C[i, h_i(s)] \geq c_s + \epsilon \|\mathbf{c}\|_1) = P(C[i, h_i(s)] - c_s \geq \epsilon \|\mathbf{c}\|_1)$$

$$\begin{array}{l} \text{(by Markov's} \\ \text{inequality)} \end{array} \leq \frac{\mathbb{E}[C[i, h_i(s)]] - c_s}{\epsilon \|\mathbf{c}\|_1}$$

$$\begin{array}{l} \text{(by bound on} \\ \mathbb{E}[C[i, h_i(s)]] \end{array} \leq \frac{c_s + \frac{\epsilon}{2} \|\mathbf{c}\|_1 - c_s}{\epsilon \|\mathbf{c}\|_1} = \frac{1}{2}$$

# Count-min Sketch: Theoretical Guarantees Proof (Cont.)

- Given  $r > \log_2 1/\delta$  hash functions that each map to  $m > 2/\epsilon$  buckets, then  $\forall s$

$$P(\hat{c}_s \geq c_s + \epsilon \|\mathbf{c}\|_1) \leq \delta$$

where  $\|\mathbf{c}\|_1$  is the total number of items in the stream.

- Proof:
  - For an arbitrary item,  $s$ :

$$\begin{aligned} P(\hat{c}_s \geq c_s + \epsilon \|\mathbf{c}\|_1) &= P\left(\min_i C[i, h_i(s)] \geq c_s + \epsilon \|\mathbf{c}\|_1\right) \\ &= P\left(\bigcap_i C[i, h_i(s)] \geq c_s + \epsilon \|\mathbf{c}\|_1\right) \end{aligned}$$

$$\text{(by independence)} \quad = \prod_i P(C[i, h_i(s)] \geq c_s + \epsilon \|\mathbf{c}\|_1)$$

$$\text{(by definition of } \delta) \quad \leq 1/2^r = \delta$$

# Count-min Sketch: Summary

- Key idea: use multiple (independent) hash functions to negate the effect of collisions in any single hash function
- Nice theoretical guarantees
- Can address overestimation bias to make unbiased count estimators (see HW3)
- Can be extended to handle weighted updates as in SGD logistic regression → efficient model updates in high-dimensional settings (large  $k$ )
  - For more details see:  
<https://arxiv.org/pdf/1711.02305.pdf>

# Two Specific Applications of Feature Hashing



1. *Count-min sketch* for logistic regression
2. *Locality sensitive hashing* for nearest neighbors/clustering

# Two Specific Applications of Feature Hashing

1. *Count-min sketch* for logistic regression
2. *Locality sensitive hashing* for nearest neighbors/clustering

# Locality-Sensitive Hashing

- Key idea: use hash functions that map “similar” items to the same function
  - Plot twist: collisions are the goal!
- Enables efficient comparison between points
- In the context of large-scale machine learning, can be used for clustering or nearest neighbor models in high-dimensional space





# Locality-Sensitive Hashing: Formal Definition

- A family of hash functions  $\mathcal{F}$  that map  $\mathbb{R}^k \rightarrow \{0, \dots, m\}$  is  $(\epsilon, c\epsilon, \delta_1, \delta_2)$ -sensitive with respect to some distance metric  $d$  iff  $\forall \mathbf{x}, \mathbf{y} \in \mathbb{R}^k$ , a hash function  $h$  chosen uniformly at random from  $\mathcal{F}$  has the property that
  - if  $d(\mathbf{x}, \mathbf{y}) \leq \epsilon$ , then  $P(h(\mathbf{x}) = h(\mathbf{y})) \geq \delta_1$ 
    - “Points close to each other (according to some distance metric  $d$ ) are likely to collide”
  - if  $d(\mathbf{x}, \mathbf{y}) \geq c\epsilon$ , then  $P(h(\mathbf{x}) = h(\mathbf{y})) \leq \delta_2$ 
    - “Points far apart (according to some distance metric  $d$ ) are unlikely to collide”

# Locality-Sensitive Hashing: Example

- Random projection:
  - $m = 1$  so our hash functions map  $\mathbb{R}^k$  to  $\{0,1\}$
  - Let  $\mathcal{F}$  be the set of all linear decision boundaries:

$$\mathcal{F} = \left\{ h : h_{\mathbf{w}}(\mathbf{x}) = \begin{cases} 1 & \mathbf{w}^T \mathbf{x} \geq 0 \\ 0 & \text{otherwise} \end{cases} \right\}$$

- $d(\mathbf{x}, \mathbf{y}) = 1 - \frac{\mathbf{x}^T \mathbf{y}}{\|\mathbf{x}\|_2 \|\mathbf{y}\|_2}$ , the *cosine distance*
  - Recall that one way to define a dot product is

$$\mathbf{x}^T \mathbf{y} = \|\mathbf{x}\|_2 \|\mathbf{y}\|_2 \cos \theta \rightarrow \cos \theta = \frac{\mathbf{x}^T \mathbf{y}}{\|\mathbf{x}\|_2 \|\mathbf{y}\|_2}$$

where  $\theta$  is the angle between the vectors  $\mathbf{x}$  and  $\mathbf{y}$

# Locality-Sensitive Hashing: Example

- Random projection:
  - $m = 1$  so our hash functions map  $\mathbb{R}^k$  to  $\{0,1\}$
  - Let  $\mathcal{F}$  be the set of all linear decision boundaries:

$$\mathcal{F} = \left\{ h : h_{\mathbf{w}}(\mathbf{x}) = \begin{cases} 1 & \mathbf{w}^T \mathbf{x} \geq 0 \\ 0 & \text{otherwise} \end{cases} \right\}$$

- $d(\mathbf{x}, \mathbf{y}) = 1 - \frac{\mathbf{x}^T \mathbf{y}}{\|\mathbf{x}\|_2 \|\mathbf{y}\|_2}$ , the *cosine distance*
- For a random weight vector  $\mathbf{w}$  (e.g., weights sampled independently from a standard Gaussian)

$$P(h_{\mathbf{w}}(\mathbf{x}) = h_{\mathbf{w}}(\mathbf{y})) = 1 - \frac{\arccos 1 - d(\mathbf{x}, \mathbf{y})}{\pi}$$

- $\mathcal{F}$  is  $\left( \epsilon, c\epsilon, 1 - \frac{\arccos 1 - \epsilon}{\pi}, 1 - \frac{\arccos 1 - c\epsilon}{\pi} \right)$ -sensitive

# Locality-Sensitive Hashing: Example

- Random projection:
  - $m = 1$  so our hash functions map  $\mathbb{R}^k$  to  $\{0,1\}$
  - Let  $\mathcal{F}$  be the set of all linear decision boundaries:

$$\mathcal{F} = \left\{ h : h_{\mathbf{w}}(\mathbf{x}) = \begin{cases} 1 & \mathbf{w}^T \mathbf{x} \geq 0 \\ 0 & \text{otherwise} \end{cases} \right\}$$


- $d(\mathbf{x}, \mathbf{y}) = 1 - \frac{\mathbf{x}^T \mathbf{y}}{\|\mathbf{x}\|_2 \|\mathbf{y}\|_2}$ , the *cosine distance*
- For a random weight vector  $\mathbf{w}$  (e.g., weights sampled independently from a standard Gaussian)

$$P(h_{\mathbf{w}}(\mathbf{x}) = h_{\mathbf{w}}(\mathbf{y})) = 1 - \frac{\theta}{\pi}$$

- $\mathcal{F}$  is  $\left( \epsilon, c\epsilon, 1 - \frac{\arccos 1-\epsilon}{\pi}, 1 - \frac{\arccos 1-c\epsilon}{\pi} \right)$ -sensitive


# Locality-Sensitive Hashing: Example

- Random projection:




What does the decision boundary corresponding to  $w$  look like?

- Random projection:




# Locality-Sensitive Hashing: Example

- Random projection:  
 $h_w(\mathbf{x}) = h_w(\mathbf{y}) = 1$



# Locality-Sensitive Hashing: Example


- Random projection:  
 $h_w(\mathbf{x}) = h_w(\mathbf{y}) = 1$






# Locality-Sensitive Hashing: Example

- Random projection:  
 $h_w(\mathbf{x}) = h_w(\mathbf{y}) = 0$




# Locality-Sensitive Hashing: Example

- Random projection:  
 $h_w(x) \neq h_w(y)$




# Locality-Sensitive Hashing: Example

- Random projection:  
 $h_w(x) \neq h_w(y)$




# Locality-Sensitive Hashing: Example

- Random projection:  
 $h_w(\mathbf{x}) \neq h_w(\mathbf{y})$



But why stop  
at just one  
weight vector?

- Random projection:  
 $h_w(\mathbf{x}) \neq h_w(\mathbf{y})$



# Locality-Sensitive Hashing: Amplification

- Random projections:
  - $m = 1$  so our hash functions map  $\mathbb{R}^k$  to  $\{0,1\}$
  - Let  $\mathcal{F}$  be the set of all linear decision boundaries:

$$\mathcal{F} = \left\{ h : h_{\mathbf{w}}(\mathbf{x}) = \begin{cases} 1 & \mathbf{w}^T \mathbf{x} \geq 0 \\ 0 & \text{otherwise} \end{cases} \right\}$$

- $d(\mathbf{x}, \mathbf{y}) = 1 - \frac{\mathbf{x}^T \mathbf{y}}{\|\mathbf{x}\|_2 \|\mathbf{y}\|_2}$ , the *cosine distance*
- Sample  $r$  weight vectors  $\mathbf{w}_1, \dots, \mathbf{w}_r$  and define a new *bit vector* representation for  $\mathbf{x}$  as


$$\mathbf{x}' = [h_{\mathbf{w}_1}(\mathbf{x}), h_{\mathbf{w}_2}(\mathbf{x}), \dots, h_{\mathbf{w}_r}(\mathbf{x})]$$

# Locality-Sensitive Hashing: Amplification

- Random projections:

$$x' = [1, 0, 1, 1, 0]$$

$$y' = [1, 0, 0, 1, 0]$$



# Locality-Sensitive Hashing: Amplification

- Random projections:
  - $m = 1$  so our hash functions map  $\mathbb{R}^k$  to  $\{0,1\}$
  - Let  $\mathcal{F}$  be the set of all linear decision boundaries:

$$\mathcal{F} = \left\{ h : h_{\mathbf{w}}(\mathbf{x}) = \begin{cases} 1 & \mathbf{w}^T \mathbf{x} \geq 0 \\ 0 & \text{otherwise} \end{cases} \right\}$$

- $d(\mathbf{x}, \mathbf{y}) = 1 - \frac{\mathbf{x}^T \mathbf{y}}{\|\mathbf{x}\|_2 \|\mathbf{y}\|_2}$ , the *cosine distance*
- Sample  $r$  weight vectors  $\mathbf{w}_1, \dots, \mathbf{w}_r$  and define a new *bit vector* representation for  $\mathbf{x}$  as

$$\mathbf{x}' = [h_{\mathbf{w}_1}(\mathbf{x}), h_{\mathbf{w}_2}(\mathbf{x}), \dots, h_{\mathbf{w}_r}(\mathbf{x})]$$

- The Hamming distance between two bit vectors is

$$d_h(\mathbf{x}', \mathbf{y}') = \sum_{i=1}^r \mathbb{1}(x'_i \neq y'_i) = \|\mathbf{x}' - \mathbf{y}'\|_1$$




# Locality-Sensitive Hashing: Amplification

- Random projections:

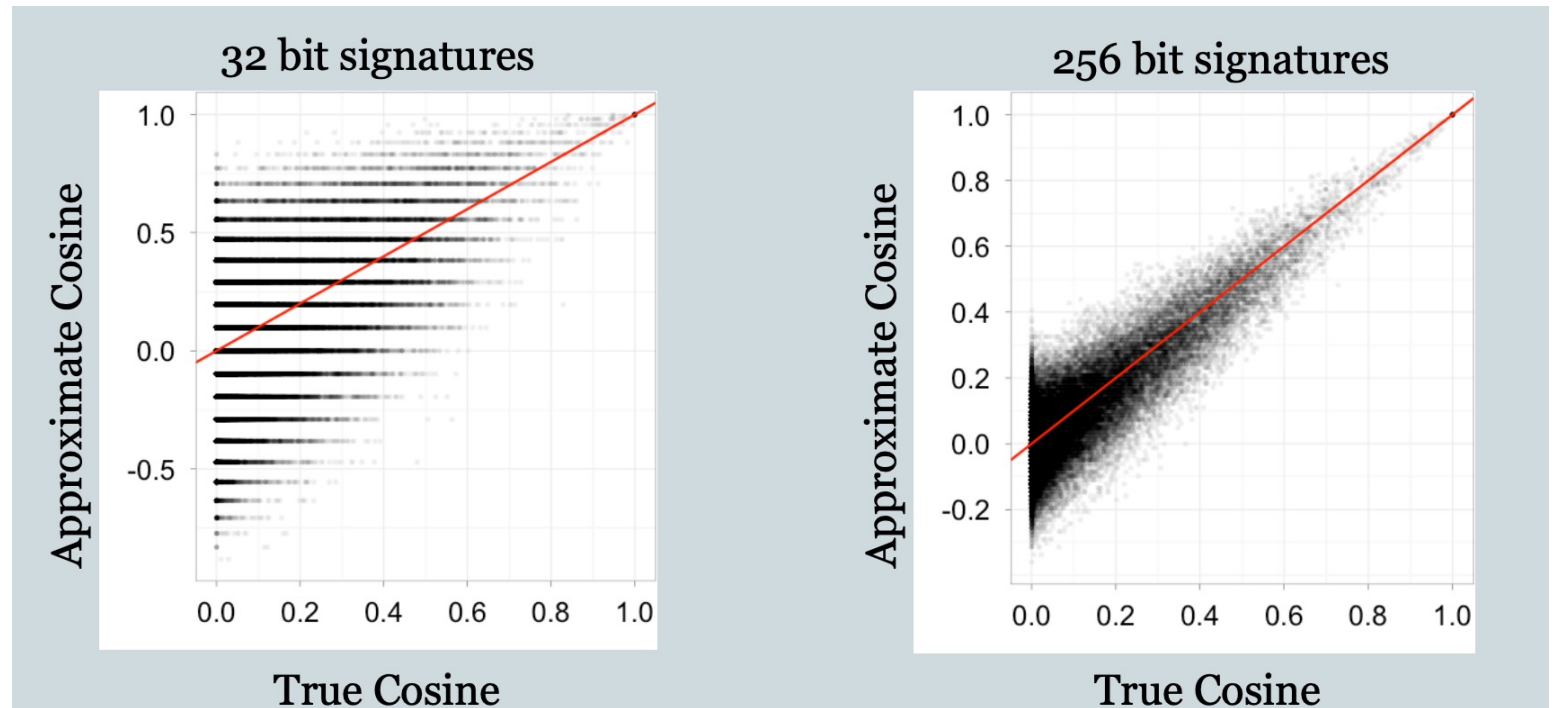
$$\mathbf{x}' = [1, 0, 1, 1, 0]$$

$$\mathbf{y}' = [1, 0, 0, 1, 0]$$



$$d(\mathbf{x}, \mathbf{y}) = 1 - \cos(\theta)$$
$$\approx 1 - \cos\left(\frac{d_h(\mathbf{x}', \mathbf{y}')}{r} \pi\right)$$

# Locality-Sensitive Hashing: Evaluation



Cheaper, less accurate approximation

Expensive, more accurate approximation

# Locality-Sensitive Hashing: Summary

- Key idea: use hash functions that map “similar” items to the same function
  - Plot twist: collisions are the goal!
- Random projections can be used to efficiently approximate the cosine distance between points
  - Reduces to computing the Hamming distance between bit vectors consisting of multiple hash functions
- Can be used to efficiently cluster data and compute the nearest neighbors to a query point

# Distributed $k$ -NN

1. Distribute training dataset across some number of worker machines
2. Communicate query point  $x$  to each worker machine
3. Map: compute distance between  $x$  and each training data point
  - If the data is high-dimensional, first compute random projection bit vectors to approximate distances
4. Reduce: return the  $k$  nearest neighbors from each worker (and their distances)
5. Driver computes the global  $k$  nearest neighbors

# Key Takeaways

- Non-numeric data can be a challenge for many machine learning models
- One-hot encodings are impractical for large datasets
  - Feature hashing is a reasonable alternative
  - Hash kernels are unbiased estimators of dot products that preserve distances (with high probability)
- For algorithms where updates are weighted counts of feature values, we can use a weighted version of count-min sketch, an application of hashing with nice theoretical guarantees and good empirical performance
- For algorithms where distances between points is relevant, locality sensitive hashing can be used to efficiently approximate certain distance metrics

# Looking unreasonably far ahead...

- Let's talk about HW4!
- Main goals:
  - Give you all hands-on experience working with cloud computing, specifically Amazon Web Services
  - Work with an actual large dataset in Spark (the entire Million Song Dataset ~280 GB)
  - Execute an end-to-end machine learning pipeline in a distributed manner
- Warning: this homework will be more open-ended than previous ones

# Looking unreasonably far ahead...

- Let's talk about HW4!
- Main goals:
  - Give you all hands-on experience working with **cloud computing**, specifically **Amazon Web Services**
  - Work with an actual large dataset in Spark (the entire Million Song Dataset ~280 GB)
  - Execute an end-to-end machine learning pipeline in a distributed manner
- Warning: this homework will be more open-ended than previous ones

## Recall: Cloud Computing

- Enables distributed processing by democratizing access to storage and computational resources
- Similar to a public utility, you can access as much or as little of it as you need



Google Cloud



## Recall: Cloud Computing

- Cloud computing relies on sharing of resources to achieve coherence and typically using a "pay-as-you-go" model which can help in reducing capital expenses but may also lead to unexpected operating expenses for users.



Google Cloud

# Data Centers: Amazon Data Center in Ashburn, Virginia



# Data Centers: Facebook Data Center in Prineville, Oregon





# Data Centers: Amazon Data Center in Brétigny-sur- Orge, France



# Cloud Computing: Pros and Cons

- Pros:
  - On-demand computing/storage allows for scalability and elasticity (only pay for what you use)
  - Eliminates infrastructure and **maintenance** costs
- Cons:
  - Depending on usage, overall (monetary) costs may be higher than being “on-perm”, i.e., buying your own compute resources
  - Data gravity – difficult to switch providers once a lot of resources have been committed to a single one
  - Data privacy – do you really trust Amazon?
  - Environmental impacts – densely packed computing resources require a lot of energy to run and **cool**

# Cloud Computing: Maintenance

## Florida Data Centers Brace for Powerful Hurricane Ian

BY RICH MILLER - SEPTEMBER 28, 2022 — [LEAVE A COMMENT](#)



## Data Center Fire: Google Suffers 'Electrical Incident,' 3 Injured

News of yet another data center fire in just two years: OVH in 2021 and now Google. Here's how to make sure you're not next.

## Apple Recovers From Massive Outage

TECH • FACEBOOK

Yet another example  
rely heavily on the cl

## Facebook's outage cost the company nearly \$100 million in revenue

BY CHRIS MORRIS

Jeff Baumgartner | Mar 22, 2021 4:30 PM EDT

Updated October 5, 2021 11:25 AM EDT

Source: <https://www.datacenterknowledge.com/google-alphabet/data-center-fire-google-suffers-electrical-incident-3-injured>


Source: <https://www.yahoo.com/now/facebook-outage-already-cost-company-203040137.html>

# Cloud Computing: Pros and Cons

- Pros:
  - On-demand computing/storage allows for scalability and elasticity (only pay for what you use)
  - Eliminates infrastructure and **maintenance** costs
- Cons:
  - Depending on usage, overall (monetary) costs may be higher than being “on-perm”, i.e., buying your own compute resources
  - Data gravity – difficult to switch providers once a lot of resources have been committed to a single one
  - Data privacy – do you really trust Amazon?
  - Environmental impacts – densely packed computing resources require a lot of energy to run and **cool**

# Cloud Computing: Services

- Cloud computing services can be thought of as different levels of abstraction
  - Software as a service (SaaS)
  - Platform as a service (PaaS)
  - Infrastructure as a service (IaaS)

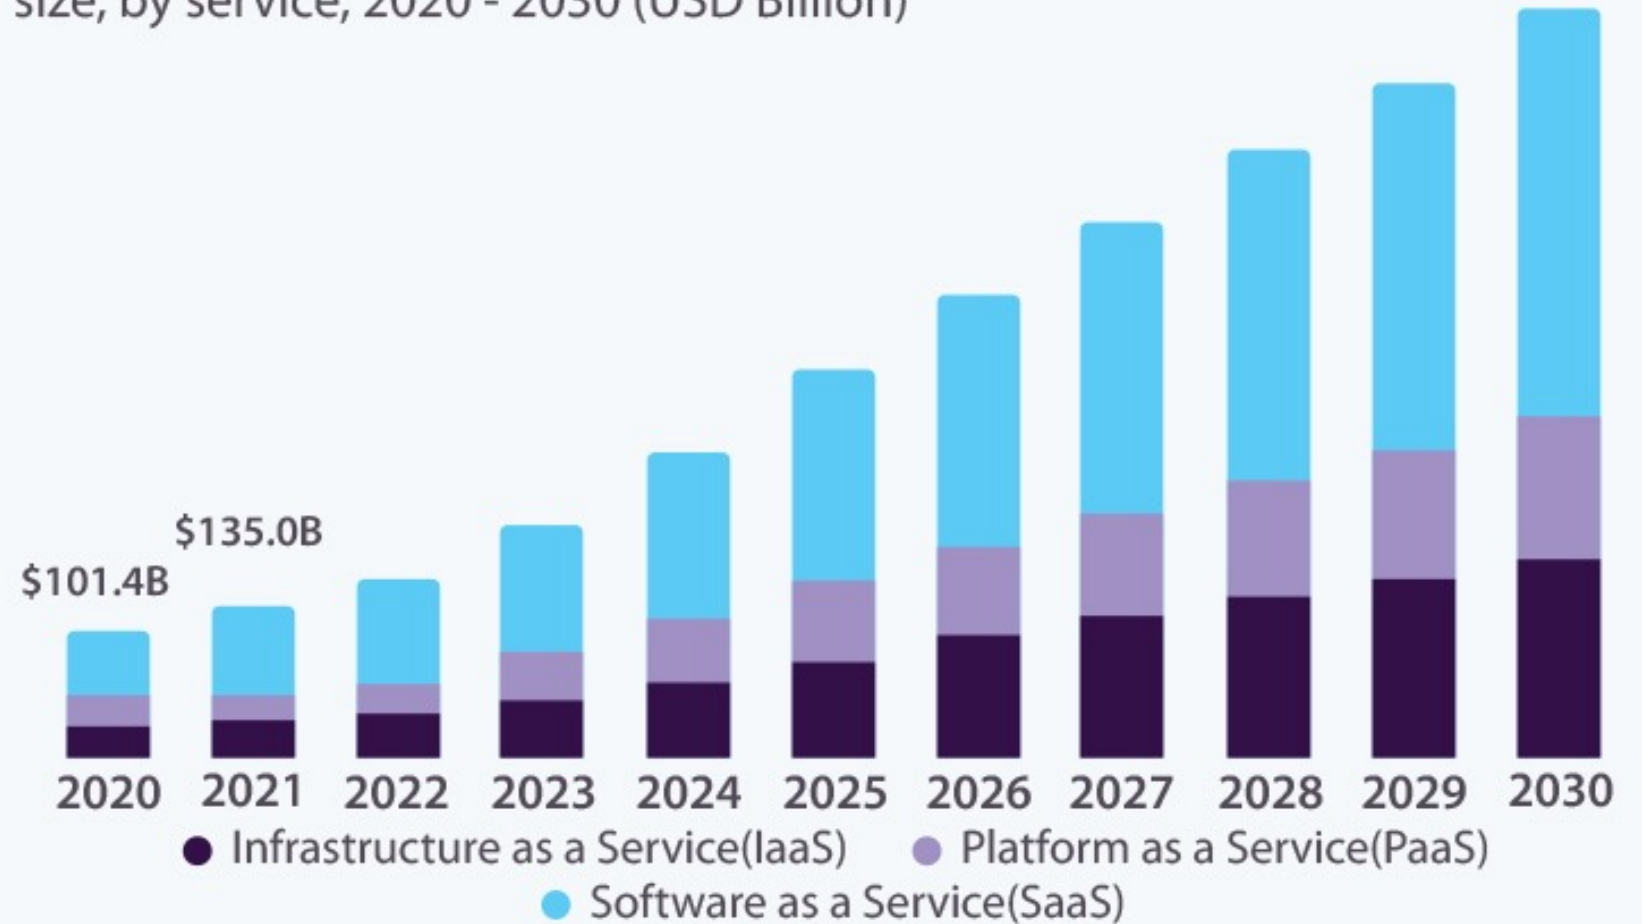




# Cloud Computing: Market

## U.S. Cloud Computing Market

size, by service, 2020 - 2030 (USD Billion)



# Cloud Computing: Market

