10417-617
Deep Learning: Fall 2020

Andrej Risteski

Machine Learning Department

Lecture 10:
Intro to unsupervised learning

Unsupervised learning

- "Pure" Reinforcement Learning (cherry)
  - The machine predicts a scalar reward given once in a while.
  - **A few bits for some samples**

- Supervised Learning (icing)
  - The machine predicts a category or a few numbers for each input
  - Predicting human-supplied data
  - **10→10,000 bits per sample**

- Unsupervised/Predictive Learning (cake)
  - The machine predicts any part of its input for any observed part.
  - Predicts future frames in videos
  - **Millions of bits per sample**

- (Yes, I know, this picture is slightly offensive to RL folks. But I'll make it up)

Slide by Yann LeCun

Unsupervised learning

Learning from data **without** labels.

What can we hope to do:

Task A: Fit a parametrized **structure** (e.g. clustering, low-dimensional subspace, manifold) to data to reveal something meaningful about data. (**Structure learning**)

Task B: Learn a (parametrized) **distribution** *close* to data generating distribution. (**Distribution learning**)

Task C: Learn a (parametrized) distribution that implicitly reveals an **"embedding"/"representation"** of data for downstream tasks. (**Representation/feature learning**)

Entangled! The "structure" and "distribution" often reveals an embedding.

Structure learning

Fit a parametrized **structure** (e.g. clustering, low-dimensional subspace) to data to reveal something meaningful about data.

PCA(principal component analysis), direction of highest variance

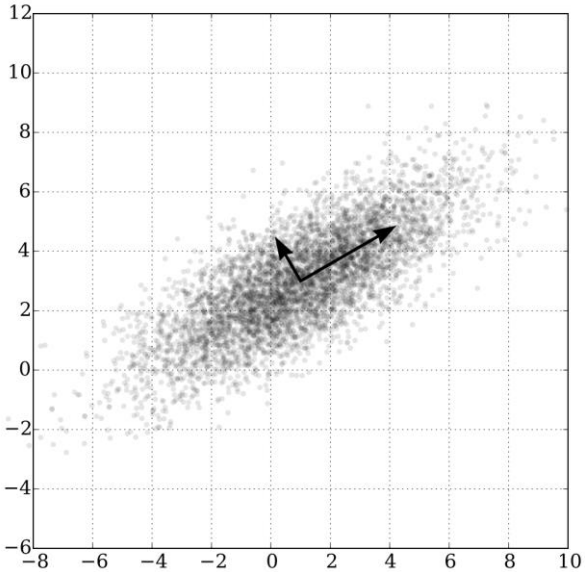

Clustering

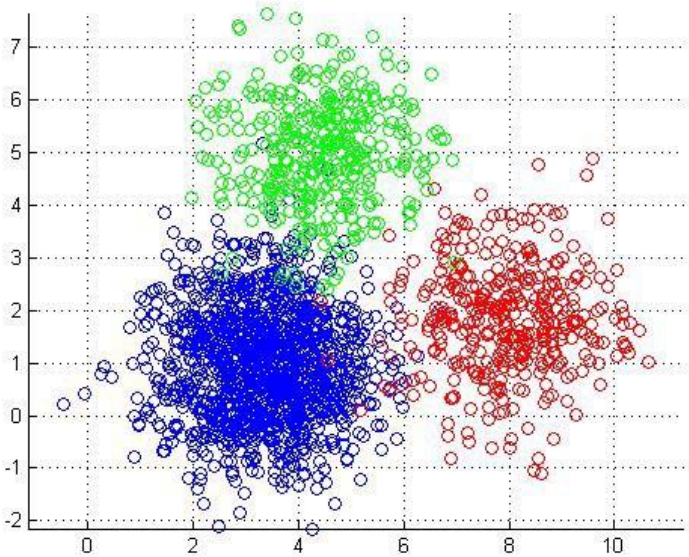

Structure learning

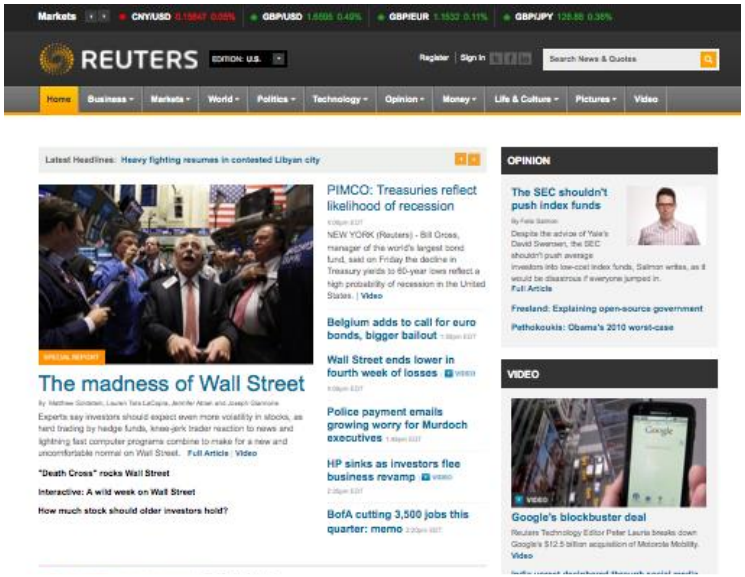

804,414 newswire stories

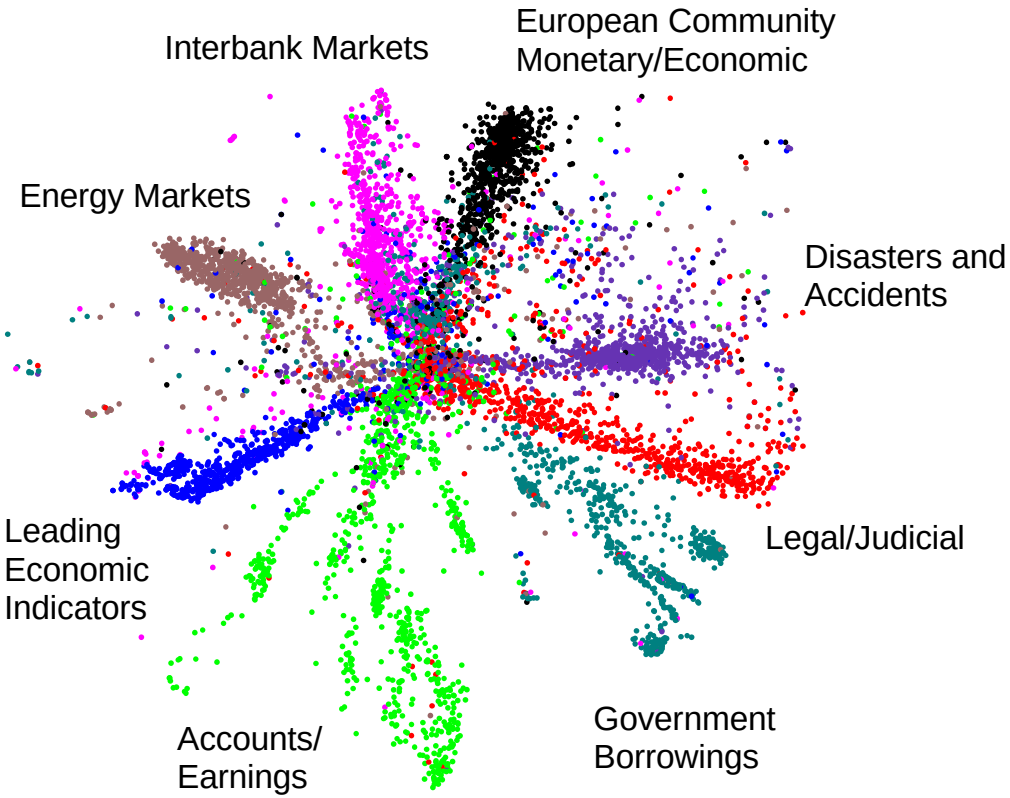

Distribution learning

Some typical choices of parametrized distributions:

Classical choices: fully-observed graphical models (undirected and directed), latent-variable graphical models (mixture models, sparse coding, topic models).

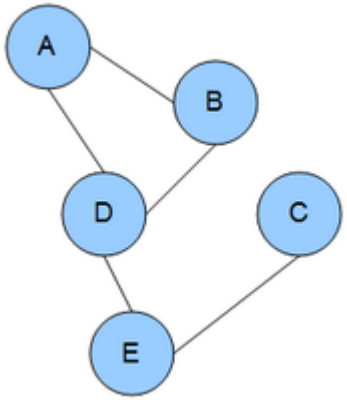

Markov Random Fields: sparse independence structure: "A is independent of other vars, given B, D"

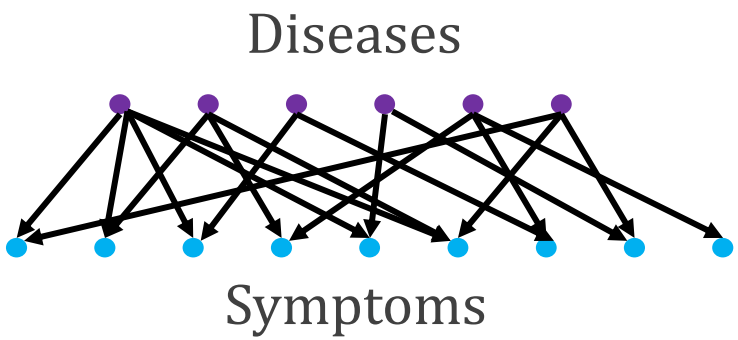

Latent variable models: data is "simple" conditioned on some unobserved (latent) variables

Distribution learning

Some typical choices of parametrized distributions:

Semi-modern choices: deep Boltzmann machines, deep belief networks, (variational) auto-encoders, energy models.

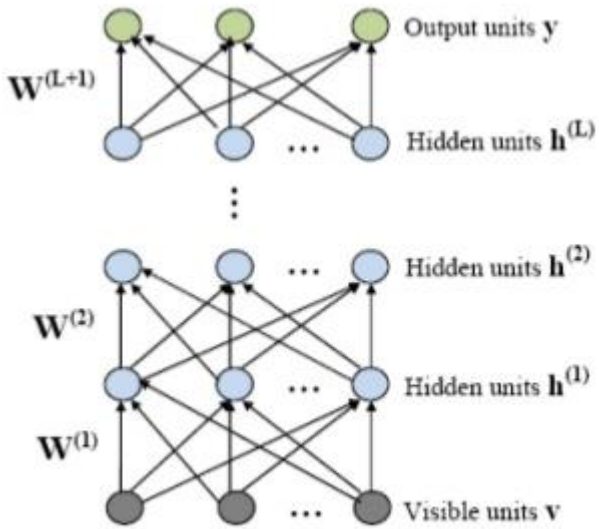

Deep Boltzmann machines, belief networks: graphical model analogues of deep neural networks.

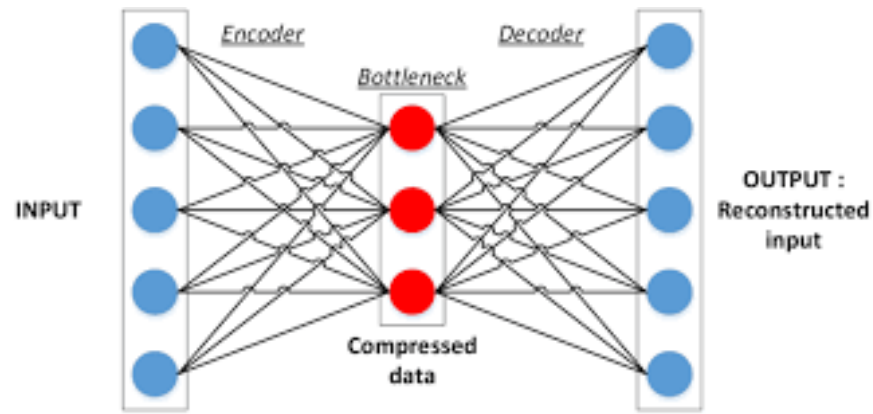

(Variational) autoencoders: model data by enforcing a latent space "bottleneck:

Distribution learning

Some typical choices of parametrized distributions:

Modern choices: generative adversarial networks, autoregressive models (pixelRNN, pixelCNN), flow models, etc.

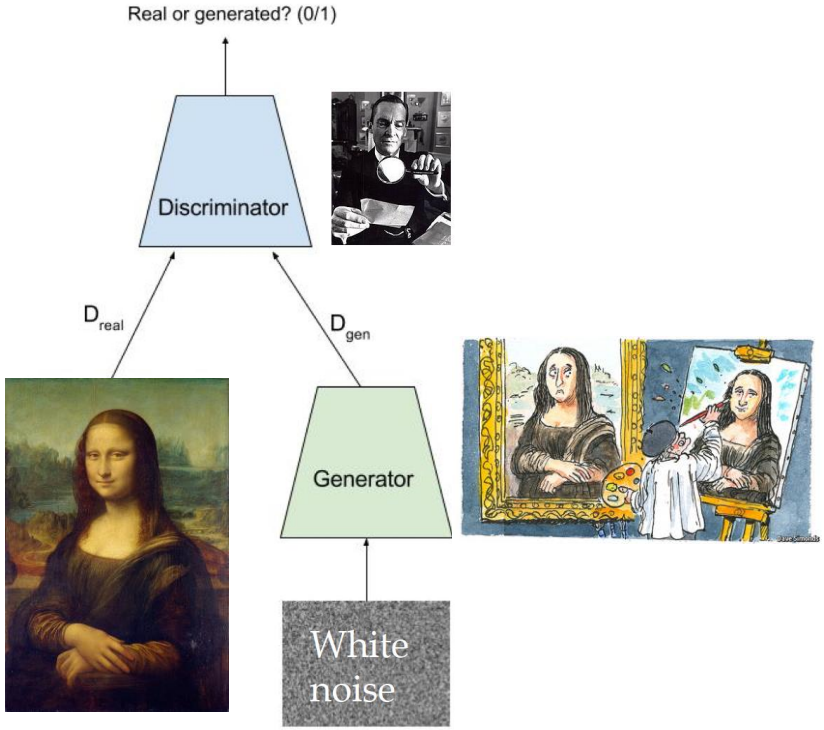

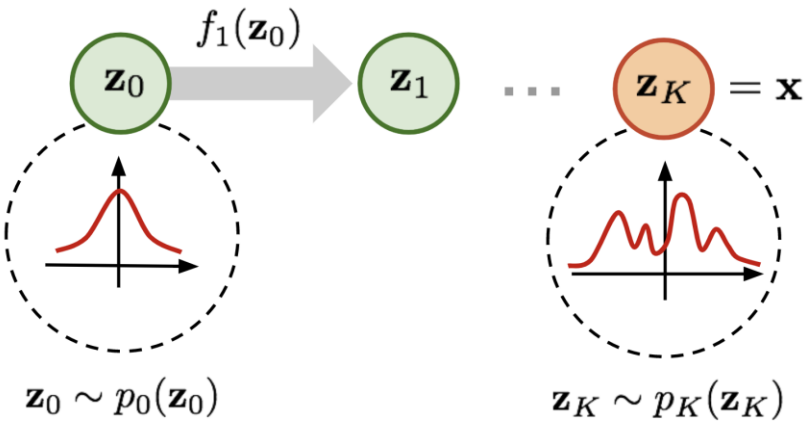

Distribution learning

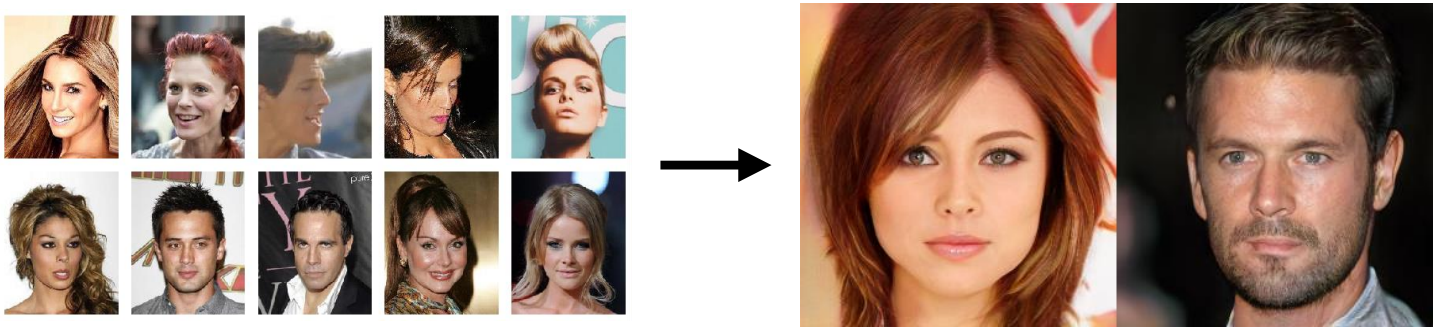

Training Data(CelebA)

Model Samples (Karras et.al., 2018)

4 years of progression on Faces

Brundage et al., 2017

Distribution learning

Whichfaceisreal.com

Distribution learning

BigGAN, Brock et al '18

Distribution learning

Conditional generative model P(zebra images| horse images)

Style Transfer

Input Image

Monet

Van Gogh

Zhou el al., Cycle GAN 2017

Distribution learning

Source actor

Real-time reenactment

Target actor

[Face2Face, Thies et al, CVPR 2016]

Representation learning and self-supervised learning

Given unlabeled data, design supervised tasks that induce a good representation for downstream tasks.

No good mathematical formalization, but the intuition is to "force" the predictor used in the task to learn something "semantically meaningful" about the data.

Examples in NLP: predict next word, given previous 5 words; predict middle word, given surrounding 5 words; etc.

Representation learning and self-supervised learning

Examples in vision: a lot, and quite different in nature.

Rotation prediction

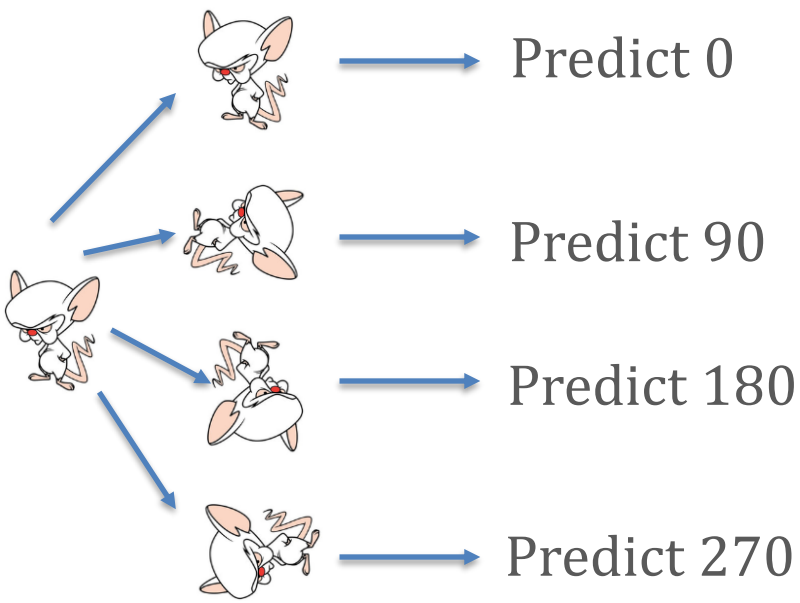

Predict one of four angles an image is rotated by

Jigsaw puzzles

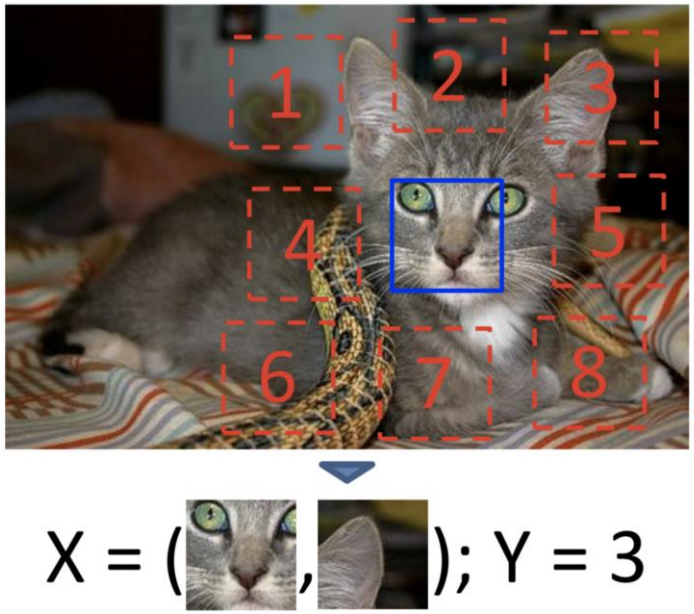

Predict position of second piece wrt to first

Relationships between the tasks

Structure learning and feature learning often blend
E.g. low dimensional features in PCA or cluster a point belongs to in clustering can be viewed as features.

Structure learning/feature learning is in general weaker than distribution learning:
E.g. methods like PCA/clustering can't be used to generate new samples.

Feature, not a bug: it has been argued that self-supervised learning works because the task we are solving is easier (both computationally and statistically)

Relationships between the tasks

Distribution learning often implies representation learning:
Many distributions we fit are **latent-variable** models (i.e. model the joint distribution between some latent variables **h** and the observed data **x**)

$$P_\theta(x, h) = P_\theta(h) P_\theta(x|h)$$

The latent variables are often viewed as a "***representation***".

The **posterior** distribution $P_\theta(h|x)$ captures a distribution over representations, given some values of the observed data.

However, a-priori, distribution $P_\theta(h|x)$ is a-priori not an easy distribution to **approximate/sample** from!

$$P_\theta(h|x) = \frac{P_\theta(x|h)p(h)}{\int_{h'} p(h')p(x|h')}$$

Hard high-dimensional sum/integral

Training techniques: Likelihood-based vs likelihood-free

Two typical families of training algorithms:

Likelihood-based: maximize the likelihood of the data under the model (possibly with some approximations)

$$\max_{\theta} \sum_{\text{samples } x_i} \log p_\theta(x_i)$$

In the limit of infinite number of samples:

$$\max_{\theta} \mathbb{E}_{x\sim p} \log p_\theta(x) = -(\mathbb{E}_{x\sim p} \log \frac{1}{p_\theta(x)} + \mathbb{E}_{x\sim p} \log p(x) - \mathbb{E}_{x\sim p} \log p(x))$$

$$= -(\quad KL(p||p_\theta) \quad + \quad H(p))$$

Hence, we are minimizing KL divergence ($H(p)$ is a constant).

Training techniques: Likelihood-based vs likelihood-free

Two typical families of training algorithms:

Likelihood-based: maximize the likelihood of the data under the model (possibly with some approximations)

Pros

Easy training: can just maximize via gradient descent.

Evaluation: evaluating the fit of the model can be done by evaluating the likelihood (on test data)

Cons

Larger models needed: likelihood objective is hard, to fit well need very big model

Likelihood encourages averaging: produced samples tend to be blurrier, as likelihood encourages "coverage" of training data.

Training techniques: Likelihood-based vs likelihood-free

$$\mathrm{KL}(p||q) = -\int p(\mathbf{Z}) \ln \frac{q(\mathbf{Z})}{p(\mathbf{Z})} \mathrm{d}\mathbf{Z}.$$

There is a large positive contribution to the KL divergence from regions of Z space in which:
- q(Z) is near zero,
- unless p(Z) is also close to zero.

Minimizing KL(p||q) leads to distributions q(Z) that are nonzero in regions where p(Z) is nonzero.

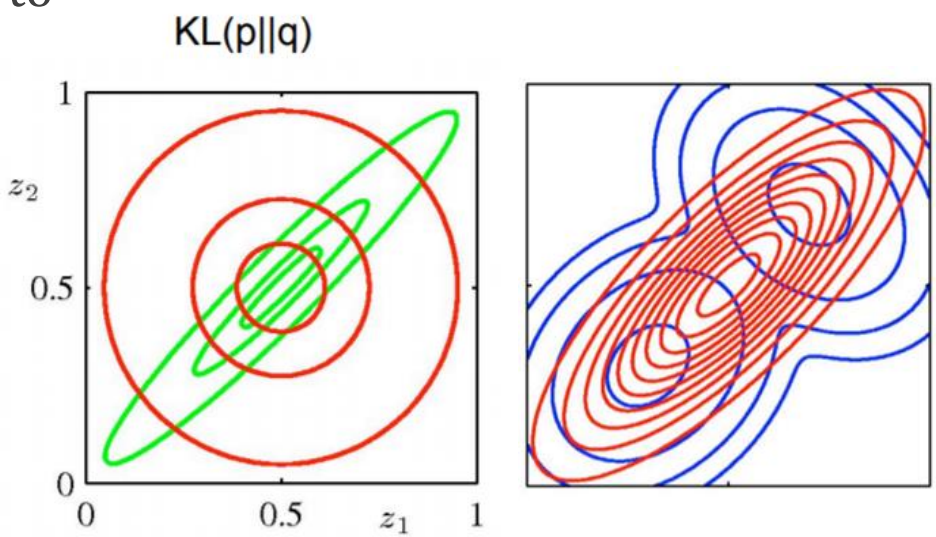

Training techniques: Likelihood-based vs likelihood-free

Faces generated using a trained VAE

Training techniques: Likelihood-based vs likelihood-free

Two typical families of training algorithms:

Likelihood-based: maximize the likelihood of the data under the model (possibly with some approximations)

Typical approximations used: variational inference (optimize tractable deterministic approximation of posteriors), MCMC methods (idea: approximate difficult quantities like posteriors with sampling)

Likelihood-free: use a surrogate loss – e.g. in GANs, train a discriminator to tell real and generated samples apart; noise-contrastive training: encourage model to put probability mass away from "fake" samples.

Training techniques: Likelihood-based vs likelihood-free

Likelihood-free: use a surrogate loss – e.g. in GANs, train a discriminator to tell real and generated samples apart; noise-contrastive training: encourage model to put probability mass away from "fake" samples.

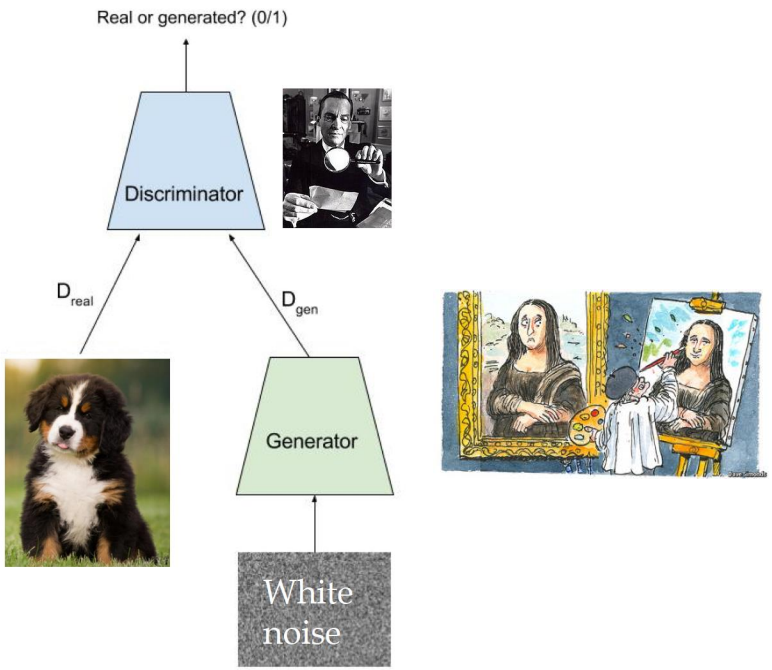

Training techniques: Likelihood-based vs likelihood-free

Likelihood-free: use a surrogate loss – e.g. in GANs, train a discriminator to tell real and generated samples apart; noise-contrastive training: encourage model to put probability mass away from "fake" samples.

Pros

Better objective, smaller models needed: objective itself (i.e. discriminator) is "learned" – can result in visually better images w/ smaller models.

Cons

Unstable training: typically min-max (saddle point) problems.

Evaluation: no way to evaluate likelihood, so no way to evaluate quality of fit.

The classics: PCA

Figure 1: Population structure within Europe.

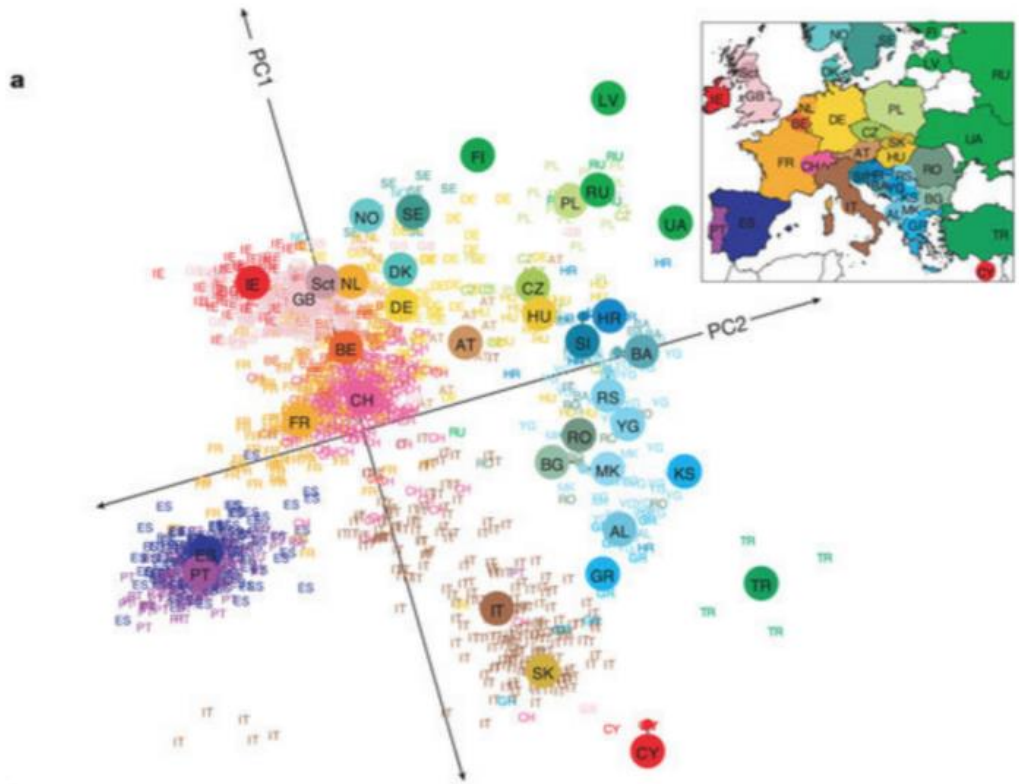

Figure 6: Plot from [1], depicting genomes for 1387 Europeans projected onto top 2 principal components. Colors/labels of datapoints correspond to geographic location of the individuals. Map of Europe (with same coloring) included in upper right for reference.

Novembre et al '08

The classics: PCA

Goal: find a k-dimensional (linear) subspace explaining most of the variance in the data.

Assume the data is centered, that is $\mathbb{E}[x] = 0$

Warmup: let k=1.

$$\underset{\{v:\,||v||=1\}}{\operatorname{argmax}} \frac{1}{m} \sum_{\text{samples } x_i} \langle v, x_i \rangle^2$$

As $m \to \infty$, we get $\mathbb{E}[\langle v, x \rangle^2]$,
i.e. variance.

The classics: PCA

Goal: find a k-dimensional (linear) subspace explaining most of the variance in the data.

$$\underset{\{v_1,v_2,\dots,v_k\}}{\operatorname{argmax}} \frac{1}{m} \sum_{\text{samples } x_i} \left(\text{length of } \mathrm{x_i} \text{ on span}(v_1, v_2, \dots, v_k)\right)^2$$

$$= \underset{\text{orthonormal } \{v_1,v_2,\dots,v_k\}}{\operatorname{argmax}} \frac{1}{m} \sum_{\text{samples } x_i} \sum_{j=1}^{k} \langle x_i, v_j \rangle^2$$

As $m \to \infty$, we get $\mathbb{E}[\sum_j \langle v_j, x \rangle^2]$

The classics: PCA

How to do this efficiently?

$$\underset{\text{orthonormal } \{v_1, v_2, \ldots, v_k\}}{\operatorname{argmax}} \frac{1}{m} \sum_{\text{samples } x_i} \sum_{j=1}^{k} \langle x_i, v_j \rangle^2$$

A convenient rewrite:

$$= \underset{\text{orthonormal } \{v_1, v_2, \ldots, v_k\}}{\operatorname{argmax}} \frac{1}{m} \sum_{\text{samples } x_i} \sum_{j=1}^{k} (v_j^T x_i)(x_i^T v_j)$$

$$= \underset{\text{orthonormal } \{v_1, v_2, \ldots, v_k\}}{\operatorname{argmax}} \frac{1}{m} \sum_{\text{samples } x_i} \sum_{j=1}^{k} v_j^T (x_i x_i^T) v_j$$

$$= \underset{\text{orthonormal } \{v_1, v_2, \ldots, v_k\}}{\operatorname{argmax}} \sum_{j=1}^{k} v_j^T \underbrace{\left(\frac{1}{m} \sum_{\text{samples } x_i} \left(x_i x_i^T \right) \right)}_{:= \text{ D, (covariance matrix)}} v_j$$

The classics: PCA

How to do this efficiently? – ***Singular Value Decomposition****!!*

$$\underset{\text{orthonormal } \{v_1, v_2, \ldots, v_k\}}{\operatorname{argmax}} \frac{1}{m} \sum_{\text{samples } x_i} \sum_{j=1}^{k} \langle x_i, v_j \rangle^2$$

A convenient rewrite: $$\underset{\text{orthonormal } \{v_1, v_2, \ldots, v_k\}}{\operatorname{argmax}} \sum_{j=1}^{k} v_j^T \mathrm{D} v_j$$

If D is **diagonal** (w/ positive entries), and we sort the entries s.t. $D_{11} \geq D_{22} \ldots \geq D_{dd}$, it's easy to see max is $\sum_{j=1}^{k} D_{jj}$. Namely:

$\sum_{j=1}^{k} v_j^T \mathrm{D} v_j = \sum_j \sum_i \left(v_j\right)_i^2 D_{ii} \leq \sum_{j=1}^{k} D_{jj}$, as $\left(v_j\right)_i^2 = 1$

The corresponding argmax is $e_1, e_2, \ldots, e_k$.

The classics: PCA

How to do this efficiently? – ***Singular Value Decomposition****!!*

$$\underset{\text{orthonormal } \{v_1, v_2, \ldots, v_k\}}{\operatorname{argmax}} \frac{1}{m} \sum_{\text{samples } x_i} \sum_{j=1}^{k} \langle x_i, v_j \rangle^2$$

A convenient rewrite: $$\underset{\text{orthonormal } \{v_1, v_2, \ldots, v_k\}}{\operatorname{argmax}} \sum_{j=1}^{k} v_j^T \mathrm{D} v_j$$

If D is not diagonal, write $\mathrm{D} = \mathrm{U}\, \widetilde{D} U^T$ to reduce to diagonal case:

$$\underset{\text{orthonormal } \{v_1, v_2, \ldots, v_k\}}{\operatorname{argmax}} \sum_{j=1}^{k} v_j^T (U\, \widetilde{D} U^T) v_j = \underset{\text{orthonormal } \{v_1, v_2, \ldots, v_k\}}{\operatorname{argmax}} \sum_{j=1}^{k} \left(U^T v_j\right)^T \widetilde{D} (U^T v_j)$$

Since U^T is orthogonal, $\{U^T v_1, U^T v_2, \ldots, U^T v_k\}$ also are orthonormal!

So, the max is $\sum_{j=1}^{k} \widetilde{D}_{jj} = \sum_{j=1}^{k} \lambda_j(D)$ and the argmax are the **top k eigenvectors** of D

The classics: clustering

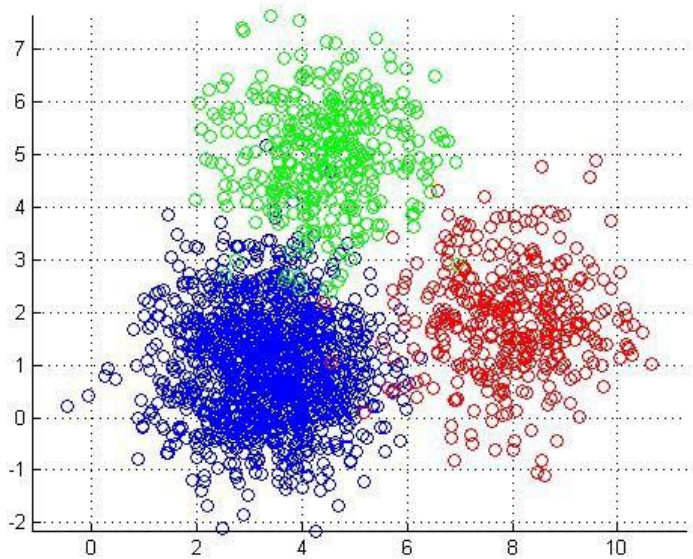

Goal: group the data into clusters of nearby points.

What's needed for clustering?

1. Proximity measure, *either*
   - similarity measure $s(x_i,x_k)$: large if x_i,x_k are similar
   - dissimilarity(or distance) measure $d(x_i,x_k)$: small if x_i,x_k are similar

2. Criterion function to evaluate a clustering

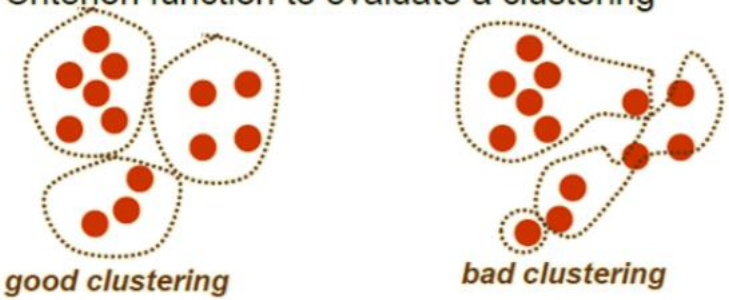

3. Algorithm to compute clustering

The classics: clustering

Popular distance metrics:

- Euclidean distance

$$d(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{\sum_{k=1}^{d} \left(x_i^{(k)} - x_j^{(k)}\right)^2}$$

  - translation invariant

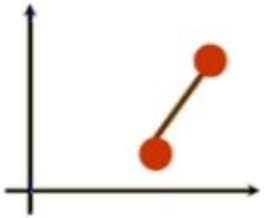

- Manhattan (city block) distance

$$d(\mathbf{x}_i, \mathbf{x}_j) = \sum_{k=1}^{d} \left|x_i^{(k)} - x_j^{(k)}\right|$$

  - approximation to Euclidean distance, cheaper to compute

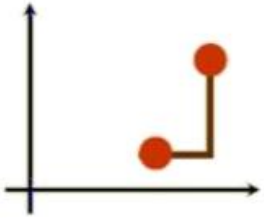

- They are special cases of **Minkowski distance**:

$$d_p(\mathbf{x}_i, \mathbf{x}_j) = \left(\sum_{k=1}^{m} \left| x_{ik} - x_{jk} \right|^p \right)^{\frac{1}{p}}$$

(p is a positive integer)

The classics: clustering

Criterion functions:

Intra-cluster cohesion
– Cohesion measures how near the data points in a cluster are to the cluster “center”.

Inter-cluster separation
– Separation means that different cluster centroids should be far away from one another.

In most applications, expert judgments are still the key

The classics: clustering

How many clusters?

- Possible approaches
  1. fix the number of clusters to ***k***
  2. find the best clustering according to the criterion function (number of clusters may vary)

K-means clustering

If the distance metric is the Euclidean distance, and the measure of cohesion is the average distance from the centroid: we get the **k-means objective.**

$$argmin_{\{r_{nk}, \boldsymbol{\mu}_\mathbf{k}\}} \sum_{n=1}^{N} \sum_{k=1}^{K} r_{nk} \|\mathbf{x_n} - \boldsymbol{\mu}_\mathbf{k}\|^2$$

Is point n in cluster k?

Centroid of k-th cluster

K-means clustering

A natural iterative algorithm, which as we will see later is a variant of the **EM (expectation-maximization)** algorithm:

Input: Data set $X = \{x^{(1)}, x^{(2)}, \cdots, x^{(m)} | x^{(i)} \in \mathbb{R}^n\}$

Output: Cluster centroids $\mu_{i=1,\cdots,k} \in \mathbb{R}^n$;Cluster assignments $c \in \mathbb{Z}$

Initialize k cluster centroids $\mu_1, \cdots, \mu_k \in \mathbb{R}^n$ randomly from X;

repeat

  for $i = 1, \cdots, m$ **do** `// Update cluster assignments`

    set $c^{(i)} = arg\ \underset{j}{min} \|x^{(i)} - \mu_j\|^2$;

  end

  for $j = 1, \cdots, k$ **do** `// Update cluster centroids`

    set $\mu_j = \frac{\sum_{i=1}^{m} 1\{c^{(i)}=j\} x^{(i)}}{\sum_{i=1}^{m} 1\{c^{(i)}=j\}}$;

  end

until *Convergence*;

return μ and c;

Algorithm 1: Algorithm of batch-version for K-means

Some weaknesses

Very sensitive to outliers:

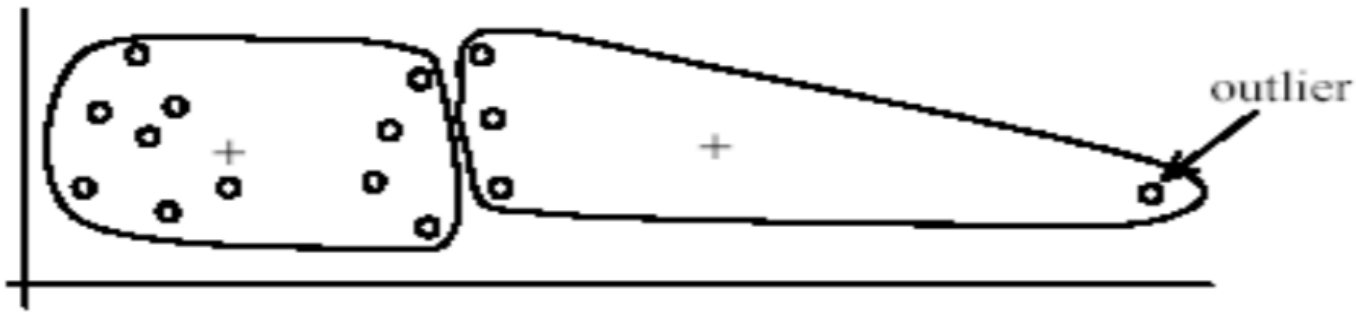

(A): Undesirable clusters

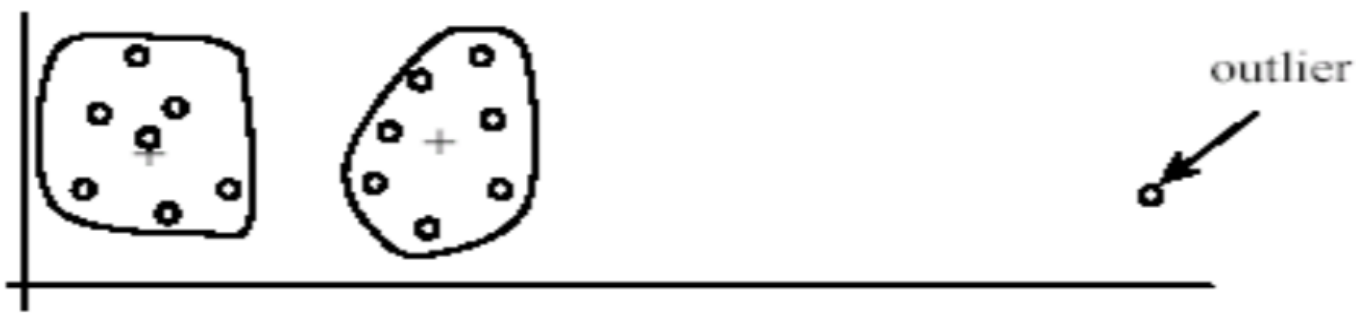

(B): Ideal clusters

Some weaknesses

Not suitable for non-spherical clusters:

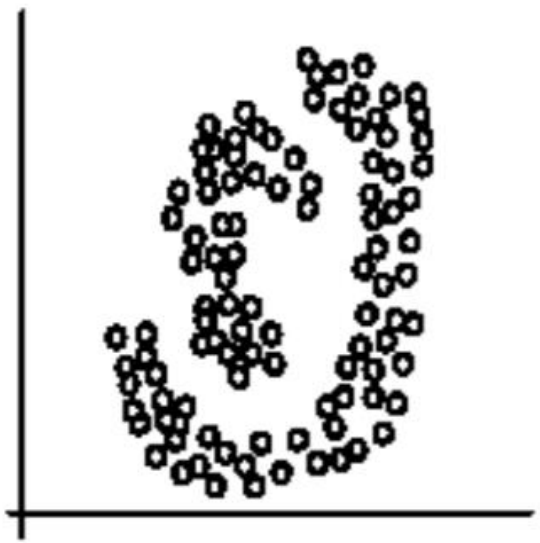

(A): Two natural clusters

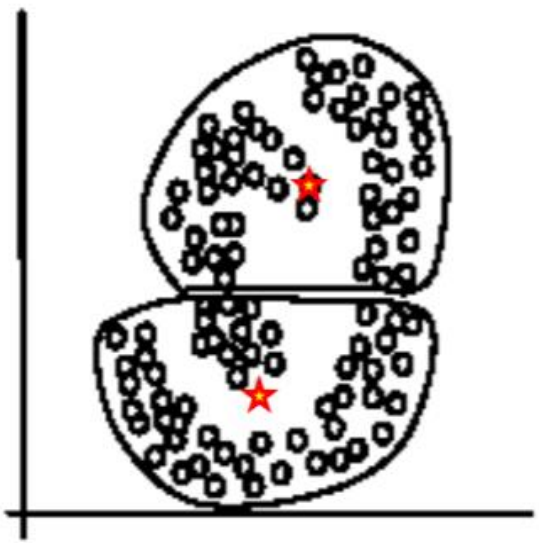

(B): *k*-means clusters