


**10417-617**  
**Deep Learning: Fall 2020**

Andrej Risteski

Machine Learning Department

**Lecture 16:**  
Generative adversarial networks,  
normalizing flows

# Some samples generated with VAEs and RBMs



Data



VAE samples

Faces generated using a trained VAE, slides from  
[http://efrosgans.eecs.berkeley.edu/CVPR18\\_slides/VAE\\_GANS\\_by\\_Rosca.pdf](http://efrosgans.eecs.berkeley.edu/CVPR18_slides/VAE_GANS_by_Rosca.pdf)

# Some samples generated with VAEs and RBMs



Faces generated using a trained VAE

# The problem

Samples are blurry, though they capture some high-level structure.


Some hypotheses for what goes wrong:

**Strong metric:** VAEs try to match the input distribution in KL divergence, which is quite a strong metric.

**Poor posteriors:** The posteriors in a VAE are Gaussian – very poor modeling power, e.g. cannot model multimodal distributions.

**Max-likelihood encourages averaging:**  
finding the max-likelihood  $q$  to fit a distribution  $p$  is equivalent to minimizing  $KL(p||q)$  (by expanding the def. of  $KL = E_p \log p - E_p \log q$ ).

Recall from when we talked about variational methods: this KL tends to “average” modes.



$KL(p||q)$




# Recap: why $\text{KL}(p||q)$ averages

$$\text{KL}(p||q) = - \int p(\mathbf{Z}) \ln \frac{q(\mathbf{Z})}{p(\mathbf{Z})} d\mathbf{Z}.$$

There is a large positive contribution to the KL divergence from regions of  $\mathbf{Z}$  space in which:

- $q(\mathbf{Z})$  is near zero,
- unless  $p(\mathbf{Z})$  is also close to zero.

Minimizing  $\text{KL}(p||q)$  leads to distributions  $q(\mathbf{Z})$  that **are nonzero in regions where  $p(\mathbf{Z})$  is nonzero.**



# The idea behind GANs

Matching a distribution on images is hard because we don't have good measures of “**distance**” between images. (Intuitively, two images could be very different in pixel space, while “**semantically**” being the same image.)

*Why don't we simultaneously train a “distance” metric as we are training the model?*

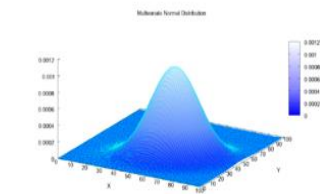
As a result, we will no longer be fitting the “maximum likelihood” model, but instead trying to learn some distribution close to the distribution of the input images in a learned metric.

This is (one of many) models which are “**likelihood-free**”: we won't be able to explicitly write a likelihood for the model, but (importantly) we will efficiently be able to draw samples from the model!


# The GAN paradigm (Goodfellow et al. '14)

**Goal:** **Learn** a distribution close to some distribution we have few samples from. (Additionally, we will be able to sample efficiently from distribution.)


**Approach:** **Fit** distribution  $P_g$  parametrized by **neural network**  $g$



$$Z \sim N(0, I_{k \times k})$$



Neural network  $g(\cdot)$



$$X = g(Z)$$

# The GAN paradigm (Goodfellow et al. '14)

Photorealistic image/video generation

MIT  
Technology  
Review

Top 10 Breakthrough  
Technologies 2018

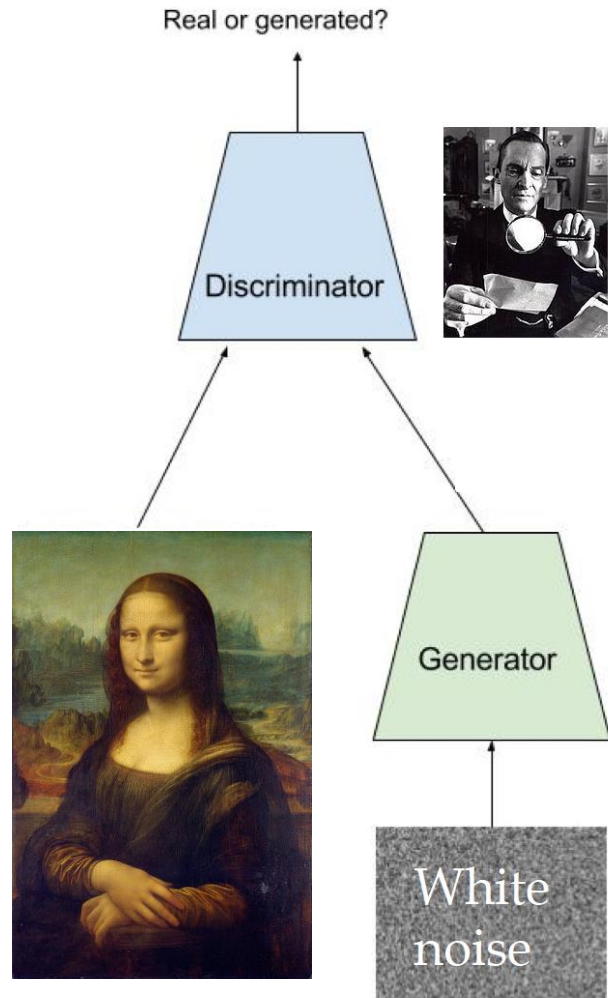


Extracting complex features





# The GAN paradigm (Goodfellow et al. '14)



Game theoretic idea:

Generator trained to **fool** discriminator.

Discriminator trained to **beat** generator.



# W-GAN formalization (Arjovsky et al. '17)

Min-max problem:

- ⊗ Min-player: generators  $g \in G$ ;    Max-player: discriminators  $f \in F$ .
- ⊗ Samples from image distr.  $P_{real}$ .    Unif. distribution over samples:  $P_{samples}$
- ⊗  $P_g$  - generator distribution:     $Z \sim N(0, I) \rightarrow g(Z)$

**Training loss:**

$$\min_{g \in G} \max_{f \in F} \left| \mathbb{E}_{P_g}[f] - \mathbb{E}_{P_{samples}}[f] \right|$$

Difference of expectation of  $f$   
on **samples vs generated**  
images



# W-GAN formalization (Arjovsky et al. '17)

Min-max problem:

- ⊗ Min-player: generators  $g \in G$ ;    Max-player: discriminators  $f \in F$ .
- ⊗ Samples from image distr.  $P_{real}$ .    Unif. distribution over samples:  $P_{samples}$
- ⊗  $P_g$  - generator distribution:     $Z \sim N(0, I) \rightarrow g(Z)$

**Training loss:**

$$\min_{g \in G} \max_{f \in F} \left| \mathbb{E}_{P_g}[f] - \mathbb{E}_{P_{samples}}[f] \right|$$

Difference of expectation of  $f$   
on **samples vs generated**  
images



# W-GAN formalization (Arjovsky et al. '17)

Min-max problem:

- ⌘ Min-player: generators  $g \in G$ ; Max-player: discriminators  $f \in F$ .
- ⌘ Samples from image distr.  $P_{real}$ . Unif. distribution over samples:  $P_{samples}$
- ⌘  $P_g$  - generator distribution:  $Z \sim N(0, I) \rightarrow g(Z)$


**Training loss:**

$$\min_{g \in G} \max_{f \in F} \left| \mathbb{E}_{P_g}[f] - \mathbb{E}_{P_{samples}}[f] \right|$$

Generator  $g$  **fools**  
discriminators  $F$  :

$$\forall f \in F, \mathbb{E}_{P_g}[f] \approx \mathbb{E}_{P_{samples}}[f]$$

Equivalently, small training  
loss!



# W-GAN formalization (Arjovsky et al. '17)


Min-max problem:

- ⊗ Min-player: generators  $g \in G$ ; Max-player: discriminators  $f \in F$ .
- ⊗ Samples from image distr.  $P_{real}$ . Unif. distribution over samples:  $P_{samples}$
- ⊗  $P_g$  - generator distribution:  $Z \sim N(0, I) \rightarrow g(Z)$

**Training loss:**

$$\min_{g \in G} \max_{f \in F} \left| \mathbb{E}_{P_g}[f] - \mathbb{E}_{P_{samples}}[f] \right|$$

Discriminators  $F$  **beat** generators if for all  $g \in G$ ,  
there is an  $f \in F$   
 $\mathbb{E}_{P_g}[f] \not\approx \mathbb{E}_{P_{samples}}[f]$





# W-GAN formalization (Arjovsky et al. '17)

Min-max problem:

- ⊗ Min-player: generators  $g \in G$ ; Max-player: discriminators  $f \in F$ .
- ⊗ Samples from image distr.  $P_{real}$ . Unif. distribution over samples:  $P_{samples}$
- ⊗  $P_g$  - generator distribution:  $Z \sim N(0, I) \rightarrow g(Z)$

**Training loss:**

$$\min_{g \in G} \max_{f \in F} \left| \mathbb{E}_{P_g}[f] - \mathbb{E}_{P_{samples}}[f] \right|$$

$$d_F(P_{samples}, P_g)$$

Discriminators  $F$  **beat** generators if for all  $g \in G$ ,  
there is an  $f \in F$   
 $\mathbb{E}_{P_g}[f] \not\approx \mathbb{E}_{P_{samples}}[f]$

“**Distance**” specified by discriminators  $F$ .  
Captures how well  $F$ ’s can **distinguish** two  
distributions

# W-GAN formalization (Arjovsky et al. '17)

Min-max problem:

- ⌘ Min-player: generators  $g \in G$ ; Max-player: discriminators  $f \in F$ .
- ⌘ Samples from image distr.  $P_{real}$ . Unif. distribution over samples:  $P_{samples}$
- ⌘  $P_g$  - generator distribution:  $Z \sim N(0, I) \rightarrow g(Z)$

**Training loss:**

$$\min_{g \in G} \max_{f \in F} \left| \mathbb{E}_{P_g}[f] - \mathbb{E}_{P_{samples}}[f] \right|$$

$$d_F(P_{samples}, P_g)$$

Discriminators  $F$  **beat** generators if for all  $g \in G$ ,  
there is an  $f \in F$   
 $\mathbb{E}_{P_g}[f] \not\approx \mathbb{E}_{P_{samples}}[f]$

$$\text{Training loss} = \min_{g \in G} d_F(P_g, P_{samples})$$

# Examples of distances $d_F$

$$\max_{f \in F} \left| \mathbb{E}_{P_g}[f] - \mathbb{E}_{P_{\text{samples}}}[f] \right|$$

$$d_F(P_{\text{samples}}, P_g)$$

Absolute value  
can be removed  
(-f is Lip if f is Lip)

$F = \{f: |f|_\infty \leq 1\}$ : **Total variation distance**

Measures differences of bounded functions

$F = \{f: \text{Lip}(f) \leq 1\}$ :  **$W_1$  (Wasserstein, earthmover) distance**

Measures differences of 1-Lipschitz functions



# Examples of distances $d_F$

$$\max_{f \in F} \left| \mathbb{E}_{P_g}[f] - \mathbb{E}_{P_{\text{samples}}}[f] \right|$$

$$d_F(P_{\text{samples}}, P_g)$$

If distance  $d_F$  is a **metric**:  $d_F(p, q) \geq 0$  and  $d_F(p, q) = 0$  only if  $p = q$

Hence, if we learn a distribution  $P_g$ , s.t.  $d_F(P_g, P_{\text{real}}) = 0$ , and  $P_{\text{real}}$  is the true data distribution, we have  $P_g = P_{\text{real}}$ .

In the limit of infinite samples,  $P_{\text{real}} = P_{\text{samples}}$ , so if training error is 0, we have learned a distribution  $P_g = P_{\text{real}}$



# Variants

Monotone function  $\phi$

$$\max_{f \in F} \mathbb{E}_{P_g}[\phi(f)] - \mathbb{E}_{P_{\text{samples}}}[\phi(1 - f)]$$

$$d_F(P_{\text{samples}}, P_g)$$

Maximize log-probability  
of correct answer

$\phi = \log$ : DC-GAN

$$\max_{f \in F} \mathbb{E}_{P_g}[\log(f)] - \mathbb{E}_{P_{\text{samples}}}[\log(1 - f)]$$

If  $F$  contains all function  $f: \mathbb{R}^d \rightarrow [0,1]$ , the above objective converges to

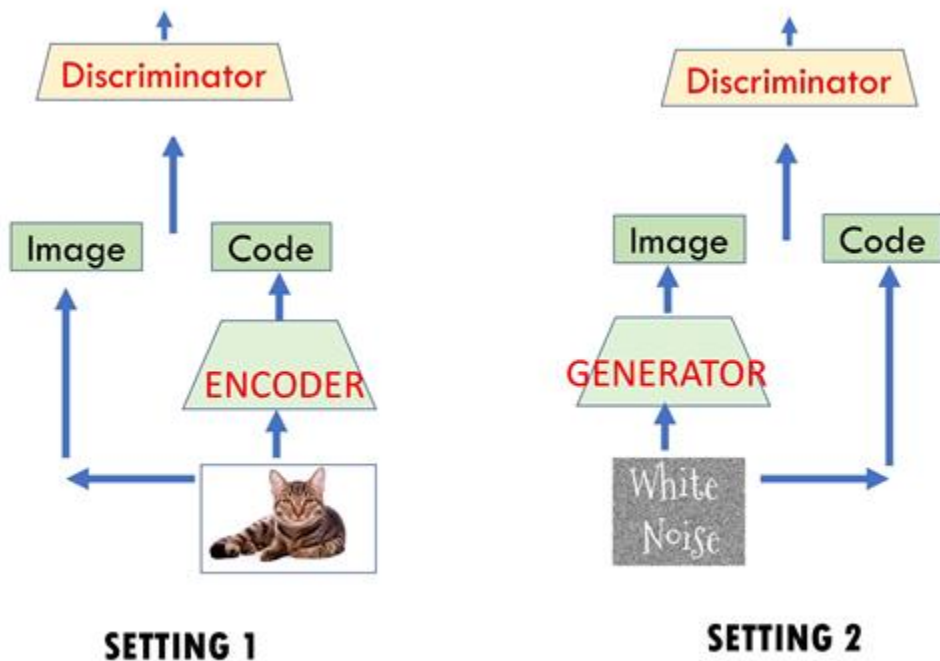
$$\min_{g \in G} \underbrace{KL(P_g || (P_{\text{real}} + P_g)/2) + KL((P_{\text{real}} + P_g)/2 || P_g)}_{\text{Jensen-Shannon divergence}} - 2 \ln 2$$

**Jensen-Shannon divergence**



# Variants

If we also want to train an **encoder E**: that is, a network that tries to output the  $z$ , s.t.  $x$  was generated from  $z$ , there is a way to adapt the adversarial setup:



Discriminator tries to distinguish b/w:

**Setting 1:** samples are  $(x, E(x))$

**Setting 2:** samples are  $(z, G(z))$

*Dumoulin et al, Donahue et al '16:*

In the limit of infinite samples, infinite capacity generators, the distributions of

Setting 1 and 2 match.

# What affects our choice of $F$ ?

**Statistical considerations:** very powerful discriminators (e.g. large neural networks) will require a lot of samples. Weak discriminators will specify a very weak metric: very “different” distributions will look very “similar” to metric.

**Our understanding here is better.**

**Algorithmic considerations:** if discriminators are very powerful, gradient information for generator is too weak and can vanish. If they are too weak – metric is weak.

**Our understanding of training dynamics is very poor.**



# Statistical questions




# Tension: strength of discriminators?

Small (weak) discriminators  $\Rightarrow$  mode collapse:

Neural net discriminators with  $\leq m$  parameters  
**fooled** by generator w/ support size  $\approx m$ .

[Arora et al'17, Arora-Risteski-Zhang ICLR'18]



Real-life distributions  
have large support!



# Tension: strength of discriminators?

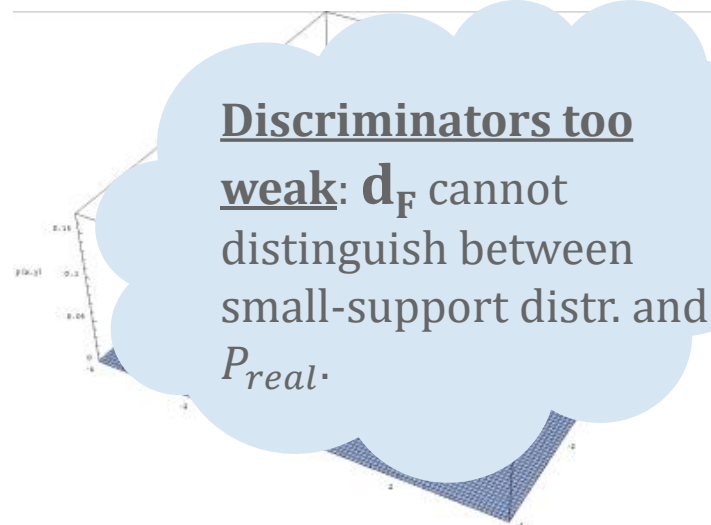
Small (weak) discriminators  $\Rightarrow$  mode collapse.

Happens for any  $P_{real}$

Neural net discriminators with  $\leq m$  parameters  
**fooled** by generator w/ support size  $\approx m$ .

[Arora et al'17, Arora-Risteski-Zhang ICLR'18]

**Not** memorization!  
More training samples  
**don't help.**



Discriminators too weak:  $d_F$  cannot distinguish between small-support distr. and  $P_{real}$ .

Real-life distributions  
have large support!



# Tension: strength of discriminators

Small discriminators  $\Rightarrow$  mode collapse:

Generator w/ support size  $\approx m$  **fools**  
neural net discriminators with  $\leq m$  **parameters**.  
[Arora et al'17, Arora-Risteski-Zhang 'ICLR18]

Large discriminators  $\Rightarrow$  poor generalization:

Loss with small # samples differs a lot from loss with infinite # samples.

$$d_F(P_{\text{samples}}, P_g) \not\approx d_F(P_{\text{real}}, P_g)$$



# Algorithmic questions



# How to train a GAN

“**Best response dynamics**”: fix generator, find best discriminator; then fix discriminator, find best generator. Repeat.

**Better in practice:** take one gradient step for generator, do a few gradient steps for discriminator. Repeat.

Going with intuition of Wasserstein distance: we'd like the discriminators to be somewhat Lipschitz – **clipping** weights is a good idea.



# How to train a GAN

How many discriminator gradient steps to take for each generator gradient step

---

**Algorithm 1** WGAN, our proposed algorithm. All experiments in the paper used the default values  $\alpha = 0.00005$ ,  $c = 0.01$ ,  $m = 64$ ,  $n_{\text{critic}} = 5$ .

---

**Require:** :  $\alpha$ , the learning rate.  $c$ , the clipping parameter.  $m$ , the batch size.  $n_{\text{critic}}$ , the number of iterations of the critic per generator iteration.

**Require:** :  $w_0$ , initial critic parameters.  $\theta_0$ , initial generator's parameters.

```
1: while  $\theta$  has not converged do
2:   for  $t = 0, \dots, n_{\text{critic}}$  do
3:     Sample  $\{x^{(i)}\}_{i=1}^m \sim \mathbb{P}_r$  a batch from the real data.
4:     Sample  $\{z^{(i)}\}_{i=1}^m \sim p(z)$  a batch of prior samples.
5:      $g_w \leftarrow \nabla_w \left[ \frac{1}{m} \sum_{i=1}^m f_w(x^{(i)}) - \frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)})) \right]$ 
6:      $w \leftarrow w + \alpha \cdot \text{RMSPProp}(w, g_w)$ 
7:      $w \leftarrow \text{clip}(w, -c, c)$ 
8:   end for
9:   Sample  $\{z^{(i)}\}_{i=1}^m \sim p(z)$  a batch of prior samples.
10:   $g_\theta \leftarrow -\nabla_\theta \frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)}))$ 
11:   $\theta \leftarrow \theta - \alpha \cdot \text{RMSPProp}(\theta, g_\theta)$ 
12: end while
```

Empirical estimates of expectations to calculate discriminator gradient

Clip

Generator gradient

Figure from Arjovsky, Chintala, Bottou '17

# Common training problems

**Unstable training:** the problem is a min-max problem (also called saddle point problem) – typically optimization is much less stable than pure minimization.

Particularly common instantiation: **cycling**




Figure 5:  $f(x,y)=(x-1/2)(y-1/2)$ .  $x$  and  $y$  are initialized at the purple diamond. Alternating between gradient ascent/ descent on  $x$  and  $y$  leads to divergent behavior, spiraling away from the optimum, but the average of the parameters is close to the optimal solution

Figure from [https://people.csail.mit.edu/madry/6.883/files/lecture\\_8.pdf](https://people.csail.mit.edu/madry/6.883/files/lecture_8.pdf)



# Common training problems

**Unstable training:** the problem is a min-max problem (also called saddle point problem) – typically optimization is much less stable than pure minimization.

**Vanishing gradient:** if the discriminator is too good, the generator gradients have a propensity to be small. (This is concerning, as to be taking gradients of the Wasserstein/JS/... objective, the discriminator needs to be optimal.)

Less of a problem with more modern GANs than with DC-GAN.




Figure 2: First, we trained a DCGAN for 1, 10 and 25 epochs. Then, with the generator fixed we train a discriminator from scratch and measure the gradients with the original cost function. We see the gradient norms decay quickly, in the best case 5 orders of magnitude after 4000 discriminator iterations. Note the logarithmic scale.

Figure from Arjovsky & Bottou '17


# Common training problems

**Unstable training:** the problem is a min-max problem (also called saddle point problem) – typically optimization is much less stable than pure minimization.

**Vanishing gradient:** if the discriminator is too good, the generator gradients have a propensity to be small. (This is concerning, as to be taking gradients of the Wasserstein/JS/... objective, the discriminator needs to be optimal.)

Less of a problem with more modern GANs than with DC-GAN.

**Mode collapse:** the training only recovers some of the modes of the underlying distribution. (**NOT** clear if this is a statistical or algorithmic problem.)




# Evaluating GANs



# How do we evaluate GANs


Wins beauty contest, but does the GAN really learn distribution?



No, just memorizing training samples.

# How do we evaluate GANs

Wins beauty contest, but does the GAN really learn distribution?



No, mode collapse: missing regions (modes) of input distribution.

# How do we evaluate GANs

Since we cannot evaluate the likelihood of the input data under a generator, **evaluation is hard**.

(Disproportionately) frequently, the evaluation is done by visually comparing samples – this cannot exclude issues like **memorization**, **mode collapse**, etc.

*Can we test for some common **failure modes**?*





# Diagnosing small support size: bday paradox



## Birthday Paradox:

If there are **23** people in a group,  $> \frac{1}{2}$  chance that two of them share a birthday.

**General version:** Suppose a distribution is uniform over  $N$  images. Then


$\Pr[\text{sample of size } \sqrt{N} \text{ has a duplicate image}] > \frac{1}{2}.$

**Birthday paradox test [Arora-Risteski-Zhang '18] :** If a sample of size  $s$  has **duplicate** images with prob.  $> \frac{1}{2}$ , then distribution essentially\* has only  **$s^2$  distinct images**.

**Implementation:** Draw sample of size  $s$ ; heuristically flag possible near-duplicates. Use human in the loop to verify duplicates.



# Diagnosing small support size: bday paradox



**CelebA (faces):** 200k training images

**DC-GAN [Radford et al.'15]:**


Support size  $\approx 250K$

**BiGAN [Donohue et al.'17] and  
ALI [Dumoulin et al.'17]:**

Support size  $\approx 1M$

A lot of **followup** and **parallel** work about diagnosing mode collapse.

# Interpolation



If linearly interpolating in latent space gives rise to meaningful images (without sharp transitions), unlikely GAN is just memorizing.

Figure from  
Radford, Metz, Chintala '16.

Figure 4: Top rows: Interpolation between a series of 9 random points in  $Z$  show that the space learned has smooth transitions, with every image in the space plausibly looking like a bedroom. In the 6th row, you see a room without a window slowly transforming into a room with a giant window. In the 10th row, you see what appears to be a TV slowly being transformed into a window.

# Inception score

Suppose we use trained network – the *Inception* architecture as a **labeler** for images. Inception gives probability over labels  $y$  for sample  $x$ :  $p(y|x)$ .

**Desirable features of generator:** the Inception classifier should be “sure” about the label for most images ( $p(y|x)$  should have *low entropy*), and the classes it generates should be diverse ( $p(y) = \mathbb{E}_{x \sim P_g} p(y|x)$  should have *high entropy*)

Thus, we want  $H(p(y|x))$  to be **low**,  $H(p(y))$  is **high**.

Consider the expression:  $\mathbb{E}_{x \sim P_g} KL(p(y|x) || p(y))$

$$= \mathbb{E}_{x \sim P_g} \mathbb{E}_{y \sim p(y|x)} \log p(y|x) - \log p(y)$$

$$= -\mathbb{E}_{x \sim P_g} H(p(y|x)) + H(p(y))$$



# Inception score

Suppose we use trained network – the *Inception* architecture as a **labeler** for images. Inception gives probability over labels  $y$  for sample  $x$ :  $p(y|x)$ .

**Desirable features of generator:** the Inception classifier should be “sure” about the label for most images ( $p(y|x)$  should have *low entropy*), and the classes it generates should be diverse ( $p(y) = \mathbb{E}_{x \sim P_g} p(y|x)$  should have *high entropy*)

Thus, we want  $H(p(y|x))$  to be **low**,  $H(p(y))$  is **high**.

Consider the expression:  $\mathbb{E}_{x \sim P_g} KL(p(y|x) || p(y))$

**Inception score:**  $\exp(\mathbb{E}_{x \sim P_g} KL(p(y|x) || p(y)))$



# Inception score

Suppose we use trained network – the *Inception* architecture as a **labeler** for images. Inception gives probability over labels  $y$  for sample  $x$ :  $p(y|x)$ .

**Desirable features of generator:** the Inception classifier should be “sure” about the label for most images ( $p(y|x)$  should have *low entropy*), and the classes it generates should be diverse ( $p(y) = \mathbb{E}_{x \sim P_g} p(y|x)$  should have *high entropy*)

Many follow ups, e.g. Frechet inception distance, modified inception score, ...

Check for **many** other metrics:

**Borji '18:** 24 quantitative, 5 qualitative measures





# The pros and cons of GANs

## *Pros*

**Photorealism:** photorealistic images, even w/ **relatively small models**.

**Efficient sampling:** easy to draw samples from model (unlike e.g. energy models).

## *Cons*

**Unstable training:** min-max problem – typically optimization much less stable than pure minimization.

**Mode collapse:** training only recovers some of the “modes” of the underlying distribution.


**Evaluation:** no likelihood, so hard to evaluate fit.



# Middle ground: “Invertible GANs/Normalizing Flows”

Can we “marry” likelihood models w/ GANs?

Suppose generator  $g: \mathbb{R}^d \rightarrow \mathbb{R}^d$  were **invertible**.



Recall from the prev. lecture: if we denote by  $\phi(z)$  the density of  $z$  under the standard Gaussian, by the change of variables formula:

$$P_g(x) = \phi(g^{-1}(x)) |\det(J_x(g^{-1}(x)))|$$

Hence, we can write down the likelihood in terms of the parameters of  $g^{-1}$  under this model!



# Middle ground:

## “Invertible GANs/Normalizing Flows”

$$P_g(x) = \phi(g^{-1}(x)) |\det(J_x(g^{-1}(x)))|$$

Hence, denoting  $g^{-1} = f_\theta$ , for some family of parametric functions  $\{f_\theta, \theta \in \Theta\}$ , the max-likelihood estimator solves

$$\max_{\theta} \sum_{i=1}^N \log \phi(f_\theta(x_i)) + \log |\det(J_x(f_\theta(x_i)))|$$


If we can evaluate and differentiate the above objective efficiently, we can do gradient-based likelihood fitting.



# Choosing invertible transforms

Note that since the change-of-variables formula composes, so if  $f_\theta = f_1 \circ f_2 \circ \dots \circ f_L$ , we have

$$\log p_\theta(x) = \sum_{i=1}^N \log \phi(f_\theta(x_i)) + \sum_{k=1}^L \log |\det(J_x(f_k(h_k(x_i))))|$$

Value of k-th layer 

So, if we can design a “simple” family of invertible transforms, we can just keep composing it.

**Try 1:** *General linear maps.*

Poor representational power: composition of linear maps is linear. If  $x = Az$ , and  $z$  is sampled from a Gaussian –  $x$  is Gaussian too.

Inefficient: Evaluating determinant of a  $d \times d$  matrix takes  $O(d^3)$  time – infeasible.



# Choosing invertible transforms

**Try 2:** *Elementwise (possibly non-linear) maps.*

Suppose that  $f_\theta(x) = (f_\theta(x_1), f_\theta(x_2), \dots, f_\theta(x_d))$

Efficient evaluation: Determinant is diagonal (since  $\frac{\partial f_\theta(x_i)}{\partial x_j} = 0$ , for  $i \neq j$ ), so

$$\det(J_x(f_\theta(x))) = \prod_i \frac{\partial f_\theta(x_i)}{\partial x_i}.$$

Poor representational power: Transforms don't “combine” coordinates.

*But, even if a matrix is triangular, Jacobian is just the product of the diagonals!!*

# Choosing invertible transforms

**Try 3:** *NICE/RealNVP (Non-linear Independent Component Estimation)*

Divide the coordinates of  $x$  into two sub-vectors with half the coords:  $x_{1:\frac{d}{2}}, x_{\frac{d}{2}+1,d}$

Divide the coordinates of  $z := f_\theta(x)$  into two sub-vectors,  $z_{1:\frac{d}{2}}, z_{\frac{d}{2}+1,d}$  and set:

$$z_{1:\frac{d}{2}} = x_{1:\frac{d}{2}}$$

$$z_{\frac{d}{2}+1,d} = x_{\frac{d}{2}+1,d} \odot \exp\left(s_\theta(x_{1:d/2})\right) + t_\theta(x_{1:d/2})$$

When is this invertible, and is the Jacobian efficiently calculated?

# Choosing invertible transforms

**Try 3:** *NICE/RealNVP (Non-linear Independent Component Estimation)*

$$z_{1:d/2} = x_{1:d/2}$$

$$z_{\frac{d}{2}+1,d} = x_{\frac{d}{2}+1,d} \odot \exp\left(s_\theta(x_{1:d/2})\right) + t_\theta(x_{1:d/2})$$

$$J_x(f_\theta(x)) = \begin{bmatrix} I & 0 \\ \frac{\partial \mathbf{z}_{d/2:d}}{\partial \mathbf{x}_{1:d/2}} & \text{diag}\left(\exp(\mathbf{s}_\theta(\mathbf{x}_{1:d/2}))\right) \end{bmatrix}$$

The determinant of a triangular matrix is the product of the diagonals!


$$\text{Hence, } \det J_x(f_\theta(x)) = \prod_i \exp\left(s_\theta(x_{1:d/2})\right)_i$$

If  $t_\theta, s_\theta$  is say, a neural net, easy to evaluate and take derivatives.




# How to choose partitions?

*NICE (Dinh et al '14), RealNVP (Dinh et al '16)*: The choice of partitions is fixed, checkerboard of channel-wise.



*Glow (Kingma et al '18)*: add trained linear transforms between affine coupling layers – i.e. a generalization of a “learned” permutation



# Some samples

## NICE/RealNVP (Non-linear Independent Component Estimation)





Figure from “*Density estimation using Real NVP*” by Dinh et al ‘16

# Some samples

*Glow*



Figure 5: Linear interpolation in latent space between real images

Figure from (Kingma et al '18)





# The pros and cons of normalizing flows

## *Pros*

**Photorealism:** photorealistic images.

**Efficient sampling:** easy to draw samples from model.

**Stabl(er) training:** likelihood objective

**Evaluation:** no likelihood, so hard to evaluate fit.

## *Cons*

**Extremely large:** in practice, good models need to be *extremely* large (Glow: 40 GPUs for ~week)

**Model depth:** training gets harder as models are typically very deep. (Glow: ~1200 layers in total)

