Computer Science > Computation and Language

arXiv:1805.04833 (cs)

[Submitted on 13 May 2018]

Hierarchical Neural Story Generation

Angela Fan, Mike Lewis, Yann Dauphin

We explore story generation: creative systems that can build coherent and fluent passages of text about a topic. We collect a large dataset of 300K human-written stories paired with writing prompts from an online forum. Our dataset enables hierarchical story generation, where the model first generates a premise, and then transforms it into a passage of text. We gain further improvements with a novel form of model fusion that improves the relevance of the story to the prompt, and adding a new gated multi-scale self-attention mechanism to model long-range context. Experiments show large improvements over strong baselines on both automated and human evaluations. Human judges prefer stories generated by our approach to those from a strong non-hierarchical model by a factor of two to one.

Subjects: **Computation and Language (cs.CL)**
Cite as: **arXiv:1805.04833 [cs.CL]**
(or **arXiv:1805.04833v1 [cs.CL]** for this version)
https://doi.org/10.48550/arXiv.1805.04833
Focus to learn more
arXiv-issued DOI via DataCite

Submission history

From: Angela Fan [view email]
[v1] Sun, 13 May 2018 07:07:08 UTC (1,135 KB)

arXivLabs: experimental projects with community collaborators

arXivLabs is a framework that allows collaborators to develop and share new arXiv features directly on our website.

Both individuals and organizations that work with arXivLabs have embraced and accepted our values of openness, community, excellence, and user data privacy. arXiv is committed to these values and only works with partners that adhere to them.

Have an idea for a project that will add value for arXiv's community? **Learn more about arXivLabs**.