

Want more content like this? **Subscribe here**

(<https://docs.google.com/forms/d/e/1FAIpQLSeOr-yp8VzYIs4ZtE9HVkRcMJyDcJ2FieM82fUsFoCsshHu9DA/viewform>) to be notified of new releases!

(<https://stanford.edu/~shervine/teaching/cs-229/cheatsheet-supervised-learning#cs-229---machine-learning>)CS 229 - Machine Learning (teaching/cs-229)

English



Supervised Learning

Unsupervised Learning

Deep Learning

Tips and tricks

# (<https://stanford.edu/~shervine/teaching/cs-229/cheatsheet-supervised-learning#cheatsheet>)Supervised Learning cheatsheet

☆ Star 15,515

By Afshine Amidi (<https://twitter.com/afshinea>) and Shervine Amidi (<https://twitter.com/shervinea>)

## (<https://stanford.edu/~shervine/teaching/cs-229/cheatsheet-supervised-learning#introduction>) Introduction to Supervised Learning

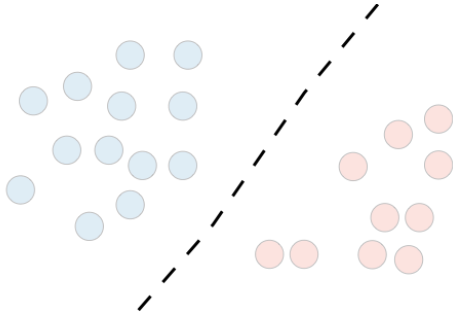
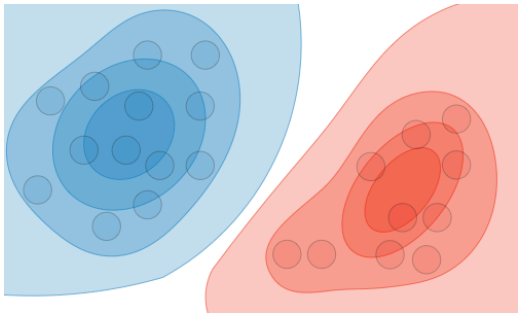
Given a set of data points  $\{x^{(1)}, \dots, x^{(m)}\}$  associated to a set of outcomes  $\{y^{(1)}, \dots, y^{(m)}\}$ , we want to build a classifier that learns how to predict  $y$  from  $x$ .

□ **Type of prediction** — The different types of predictive models are summed up in the table below:

	Regression	Classification
Outcome	Continuous	Class
Examples	Linear regression	Logistic regression, SVM, Naive Bayes

□ **Type of model** — The different models are summed up in the table below:

	Discriminative model	Generative model
--	----------------------	------------------

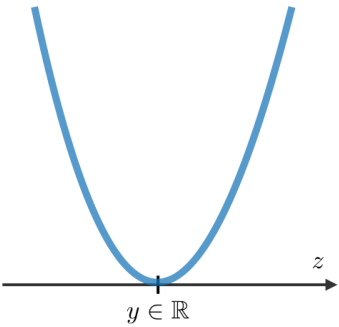
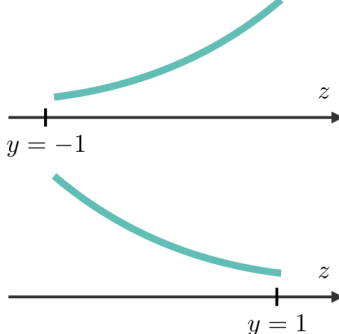
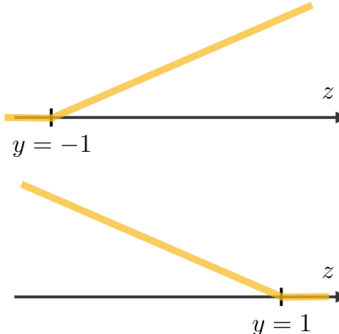
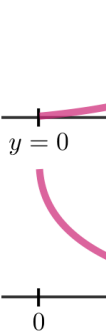
<b>Goal</b>	Directly estimate $P(y x)$	Estimate $P(x y)$ to then deduce $P(y x)$
<b>What's learned</b>	Decision boundary	Probability distributions of the data
<b>Illustration</b>		
<b>Examples</b>	Regressions, SVMs	GDA, Naive Bayes

(<https://stanford.edu/~shervine/teaching/cs-229/cheatsheet-supervised-learning#notations>)

## Notations and general concepts

□ **Hypothesis** — The hypothesis is noted  $h_\theta$  and is the model that we choose. For a given input data  $x^{(i)}$  the model prediction output is  $h_\theta(x^{(i)})$ .

□ **Loss function** — A loss function is a function  $L : (z, y) \in \mathbb{R} \times Y \mapsto L(z, y) \in \mathbb{R}$  that takes as inputs the predicted value  $z$  corresponding to the real data value  $y$  and outputs how different they are. The common loss functions are summed up in the table below:


Least squared error	Logistic loss	Hinge loss	Croc
$\frac{1}{2}(y - z)^2$	$\log(1 + \exp(-yz))$	$\max(0, 1 - yz)$	$-\left[\begin{matrix} y \\ y \end{matrix}\right]$
			
Linear regression	Logistic regression	SVM	Neu

□ **Cost function** — The cost function  $J$  is commonly used to assess the performance of a model, and is defined with the loss function  $L$  as follows:

$$J(\theta) = \sum_{i=1}^m L(h_{\theta}(x^{(i)}), y^{(i)})$$

□ **Gradient descent** — By noting  $\alpha \in \mathbb{R}$  the learning rate, the update rule for gradient descent is expressed with the learning rate and the cost function  $J$  as follows:

$$\theta \leftarrow \theta - \alpha \nabla J(\theta)$$



*Remark: Stochastic gradient descent (SGD) is updating the parameter based on each training example, and batch gradient descent is on a batch of training examples.*

□ **Likelihood** — The likelihood of a model  $L(\theta)$  given parameters  $\theta$  is used to find the optimal parameters  $\theta$  through likelihood maximization. We have:

$$\theta^{\text{opt}} = \arg \max_{\theta} L(\theta)$$

*Remark: in practice, we use the log-likelihood  $\ell(\theta) = \log(L(\theta))$  which is easier to optimize.*

□ **Newton's algorithm** — Newton's algorithm is a numerical method that finds  $\theta$  such that  $\ell'(\theta) = 0$ . Its update rule is as follows:

$$\theta \leftarrow \theta - \frac{\ell'(\theta)}{\ell''(\theta)}$$

*Remark: the multidimensional generalization, also known as the Newton-Raphson method, has the following update rule:*

$$\theta \leftarrow \theta - (\nabla_{\theta}^2 \ell(\theta))^{-1} \nabla_{\theta} \ell(\theta)$$

<https://stanford.edu/~shervine/teaching/cs-229/cheatsheet-supervised-learning#linear-models>

## Linear models

### Linear regression

We assume here that  $y|x; \theta \sim \mathcal{N}(\mu, \sigma^2)$

□ **Normal equations** — By noting  $X$  the design matrix, the value of  $\theta$  that minimizes the cost function is a closed-form solution such that:

$$\theta = (X^T X)^{-1} X^T y$$

□ **LMS algorithm** — By noting  $\alpha$  the learning rate, the update rule of the Least Mean Squares (LMS) algorithm for a training set of  $m$  data points, which is also known as the Widrow-Hoff learning rule, is as follows:

$$\forall j, \quad \theta_j \leftarrow \theta_j + \alpha \sum_{i=1}^m \left[ y^{(i)} - h_{\theta}(x^{(i)}) \right] x_j^{(i)}$$

*Remark: the update rule is a particular case of the gradient ascent.*

□ **LWR** — Locally Weighted Regression, also known as LWR, is a variant of linear regression that weights each training example in its cost function by  $w^{(i)}(x)$ , which is defined with parameter  $\tau \in \mathbb{R}$  as:

$$w^{(i)}(x) = \exp \left( -\frac{(x^{(i)} - x)^2}{2\tau^2} \right)$$

## Classification and logistic regression

□ **Sigmoid function** — The sigmoid function  $g$ , also known as the logistic function, is defined as follows:

$$\forall z \in \mathbb{R}, \quad g(z) = \frac{1}{1 + e^{-z}} \in ]0, 1[$$

□ **Logistic regression** — We assume here that  $y|x; \theta \sim \text{Bernoulli}(\phi)$ . We have the following form:

$$\phi = p(y = 1|x; \theta) = \frac{1}{1 + \exp(-\theta^T x)} = g(\theta^T x)$$

*Remark: logistic regressions do not have closed form solutions.*

□ **Softmax regression** — A softmax regression, also called a multiclass logistic regression, is used to generalize logistic regression when there are more than 2 outcome classes. By convention, we set  $\theta_K = 0$ , which makes the Bernoulli parameter  $\phi_i$  of each class  $i$  be such that:

$$\phi_i = \frac{\exp(\theta_i^T x)}{\sum_{j=1}^K \exp(\theta_j^T x)}$$

## Generalized Linear Models

□ **Exponential family** — A class of distributions is said to be in the exponential family if it can be written in terms of a natural parameter, also called the canonical parameter or link function,  $\eta$ , a sufficient statistic  $T(y)$  and a log-partition function  $a(\eta)$  as follows:

$$p(y; \eta) = b(y) \exp(\eta T(y) - a(\eta))$$

*Remark: we will often have  $T(y) = y$ . Also,  $\exp(-a(\eta))$  can be seen as a normalization parameter that will make sure that the probabilities sum to one.*

The most common exponential distributions are summed up in the following table:

Distribution	$\eta$	$T(y)$	$a(\eta)$	$b(y)$
Bernoulli	$\log\left(\frac{\phi}{1-\phi}\right)$	$y$	$\log(1 + \exp(\eta))$	1
Gaussian	$\mu$	$y$	$\frac{\eta^2}{2}$	$\frac{1}{\sqrt{2\pi}} \exp\left(-\frac{y^2}{2}\right)$
Poisson	$\log(\lambda)$	$y$	$e^\eta$	$\frac{1}{y!}$
Geometric	$\log(1 - \phi)$	$y$	$\log\left(\frac{e^\eta}{1-e^\eta}\right)$	1

□ **Assumptions of GLMs** — Generalized Linear Models (GLM) aim at predicting a random variable  $y$  as a function of  $x \in \mathbb{R}^{n+1}$  and rely on the following 3 assumptions:

- (1)  $y|x; \theta \sim \text{ExpFamily}(\eta)$
- (2)  $h_\theta(x) = E[y|x; \theta]$
- (3)  $\eta = \theta^T x$

*Remark: ordinary least squares and logistic regression are special cases of generalized linear models.*

## <https://stanford.edu/~shervine/teaching/cs-229/cheatsheet-supervised-learning#svm> Support Vector Machines


The goal of support vector machines is to find the line that maximizes the minimum distance to the line.

□ **Optimal margin classifier** — The optimal margin classifier  $h$  is such that:

$$h(x) = \text{sign}(w^T x - b)$$

where  $(w, b) \in \mathbb{R}^n \times \mathbb{R}$  is the solution of the following optimization problem:

$$\min \frac{1}{2} \|w\|^2 \quad \text{such that} \quad y^{(i)}(w^T x^{(i)} - b) \geq 1$$



Remark: the decision boundary is defined as  $w^T x - b = 0$ .


□ **Hinge loss** — The hinge loss is used in the setting of SVMs and is defined as follows:

$$L(z, y) = [1 - yz]_+ = \max(0, 1 - yz)$$

□ **Kernel** — Given a feature mapping  $\phi$ , we define the kernel  $K$  as follows:

$$K(x, z) = \phi(x)^T \phi(z)$$

In practice, the kernel  $K$  defined by  $K(x, z) = \exp\left(-\frac{\|x-z\|^2}{2\sigma^2}\right)$  is called the Gaussian kernel and is commonly used.



Non-linear separability  $\longrightarrow$  Use of a kernel mapping  $\phi$   $\longrightarrow$  Decision boundary in the original space

Remark: we say that we use the "kernel trick" to compute the cost function using the kernel because we actually don't need to know the explicit mapping  $\phi$ , which is often very complicated. Instead, only the values  $K(x, z)$  are needed.

□ **Lagrangian** — We define the Lagrangian  $\mathcal{L}(w, b)$  as follows:

$$\mathcal{L}(w, b) = f(w) + \sum_{i=1}^l \beta_i h_i(w)$$

*Remark: the coefficients  $\beta_i$  are called the Lagrange multipliers.*

## <https://stanford.edu/~shervine/teaching/cs-229/cheatsheet-supervised-learning#generative-learning> Generative Learning

A generative model first tries to learn how the data is generated by estimating  $P(x|y)$ , which we can then use to estimate  $P(y|x)$  by using Bayes' rule.

### Gaussian Discriminant Analysis

□ **Setting** — The Gaussian Discriminant Analysis assumes that  $y$  and  $x|y = 0$  and  $x|y = 1$  are such that:

- (1)  $y \sim \text{Bernoulli}(\phi)$
- (2)  $x|y = 0 \sim \mathcal{N}(\mu_0, \Sigma)$
- (3)  $x|y = 1 \sim \mathcal{N}(\mu_1, \Sigma)$

□ **Estimation** — The following table sums up the estimates that we find when maximizing the likelihood:

$\hat{\phi}$	$\hat{\mu}_j \quad (j = 0, 1)$	$\hat{\Sigma}$
$\frac{1}{m} \sum_{i=1}^m 1_{\{y^{(i)}=1\}}$	$\frac{\sum_{i=1}^m 1_{\{y^{(i)}=j\}} x^{(i)}}{\sum_{i=1}^m 1_{\{y^{(i)}=j\}}}$	$\frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu_{y^{(i)}})(x^{(i)} - \mu_{y^{(i)}})^T$

### Naive Bayes

□ **Assumption** — The Naive Bayes model supposes that the features of each data point are all independent:



$$P(x|y) = P(x_1, x_2, \dots | y) = P(x_1|y)P(x_2|y)\dots = \prod_{i=1}^n P(x_i|y)$$

□ **Solutions** — Maximizing the log-likelihood gives the following solutions:

$$P(y = k) = \frac{1}{m} \times \#\{j | y^{(j)} = k\} \quad \text{and} \quad P(x_i = l | y = k) = \frac{\#\{j | y^{(j)} = k \text{ and } x_i^{(j)} = l\}}{\#\{j | y^{(j)} = k\}}$$

with  $k \in \{0, 1\}$  and  $l \in [1, L]$

*Remark: Naive Bayes is widely used for text classification and spam detection.*

<https://stanford.edu/~shervine/teaching/cs-229/cheatsheet-supervised-learning#tree>

## Tree-based and ensemble methods

These methods can be used for both regression and classification problems.

□ **CART** — Classification and Regression Trees (CART), commonly known as decision trees, can be represented as binary trees. They have the advantage to be very interpretable.

□ **Random forest** — It is a tree-based technique that uses a high number of decision trees built out of randomly selected sets of features. Contrary to the simple decision tree, it is highly uninterpretable but its generally good performance makes it a popular algorithm.

*Remark: random forests are a type of ensemble methods.*

□ **Boosting** — The idea of boosting methods is to combine several weak learners to form a stronger one. The main ones are summed up in the table below:


Adaptive boosting	Gradient boosting
<ul style="list-style-type: none"> <li>• High weights are put on errors to improve at the next boosting step</li> <li>• Known as Adaboost</li> </ul>	<ul style="list-style-type: none"> <li>• Weak learners are trained on residuals</li> <li>• Examples include XGBoost</li> </ul>

(<https://stanford.edu/~shervine/teaching/cs-229/cheatsheet-supervised-learning#other>)

## Other non-parametric approaches

□  **$k$ -nearest neighbors** — The  $k$ -nearest neighbors algorithm, commonly known as  $k$ -NN, is a non-parametric approach where the response of a data point is determined by the nature of its  $k$  neighbors from the training set. It can be used in both classification and regression settings.

*Remark: the higher the parameter  $k$ , the higher the bias, and the lower the parameter  $k$ , the higher the variance.*




(<https://stanford.edu/~shervine/teaching/cs-229/cheatsheet-supervised-learning#learning-theory>)


## Learning Theory

□ **Union bound** — Let  $A_1, \dots, A_k$  be  $k$  events. We have:


$$P(A_1 \cup \dots \cup A_k) \leq P(A_1) + \dots + P(A_k)$$




$A_1 \cup A_2 \cup A_3$



$A_1$



$A_2$



$A_3$

□ **Hoeffding inequality** — Let  $Z_1, \dots, Z_m$  be  $m$  iid variables drawn from a Bernoulli distribution of parameter  $\phi$ . Let  $\hat{\phi}$  be their sample mean and  $\gamma > 0$  fixed. We have:

$$P(|\phi - \hat{\phi}| > \gamma) \leq 2 \exp(-2\gamma^2 m)$$

*Remark: this inequality is also known as the Chernoff bound.*

□ **Training error** — For a given classifier  $h$ , we define the training error  $\hat{\epsilon}(h)$ , also known as the empirical risk or empirical error, to be as follows:

$$\hat{\epsilon}(h) = \frac{1}{m} \sum_{i=1}^m 1_{\{h(x^{(i)}) \neq y^{(i)}\}}$$

□ **Probably Approximately Correct (PAC)** — PAC is a framework under which numerous results on learning theory were proved, and has the following set of assumptions:

- the training and testing sets follow the same distribution
- the training examples are drawn independently

□ **Shattering** — Given a set  $S = \{x^{(1)}, \dots, x^{(d)}\}$ , and a set of classifiers  $\mathcal{H}$ , we say that  $\mathcal{H}$  shatters  $S$  if for any set of labels  $\{y^{(1)}, \dots, y^{(d)}\}$ , we have:

$$\exists h \in \mathcal{H}, \quad \forall i \in \llbracket 1, d \rrbracket, \quad h(x^{(i)}) = y^{(i)}$$

□ **Upper bound theorem** — Let  $\mathcal{H}$  be a finite hypothesis class such that  $|\mathcal{H}| = k$  and let  $\delta$  and the sample size  $m$  be fixed. Then, with probability of at least  $1 - \delta$ , we have:

$$\epsilon(\hat{h}) \leq \left( \min_{h \in \mathcal{H}} \epsilon(h) \right) + 2 \sqrt{\frac{1}{2m} \log \left( \frac{2k}{\delta} \right)}$$

□ **VC dimension** — The Vapnik-Chervonenkis (VC) dimension of a given infinite hypothesis class  $\mathcal{H}$ , noted  $\text{VC}(\mathcal{H})$  is the size of the largest set that is shattered by  $\mathcal{H}$ .

*Remark: the VC dimension of  $\mathcal{H} = \{\text{set of linear classifiers in 2 dimensions}\}$  is 3.*



□ **Theorem (Vapnik)** — Let  $\mathcal{H}$  be given, with  $VC(\mathcal{H}) = d$  and  $m$  the number of training examples. With probability at least  $1 - \delta$ , we have:

$$\epsilon(\hat{h}) \leq \left( \min_{h \in \mathcal{H}} \epsilon(h) \right) + O \left( \sqrt{\frac{d}{m} \log \left( \frac{m}{d} \right) + \frac{1}{m} \log \left( \frac{1}{\delta} \right)} \right)$$



(<https://twitter.com/shervinea>)



(<https://linkedin.com/in/shervineamidi>)



(<https://github.com/shervinea>)



(<https://scholar.google.com/citations?user=nMnMTm8AAAAJ>)

