



Bias-Variance Tradeoff

CS229: Machine Learning


Carlos Guestrin

Stanford University


Slides include content developed by and co-developed with Emily Fox



Fit data with a line or ... ?




What about a quadratic function?




©2021 Carlos Guestrin


Even higher order polynomial



Do you believe this fit?




Do you believe this fit?



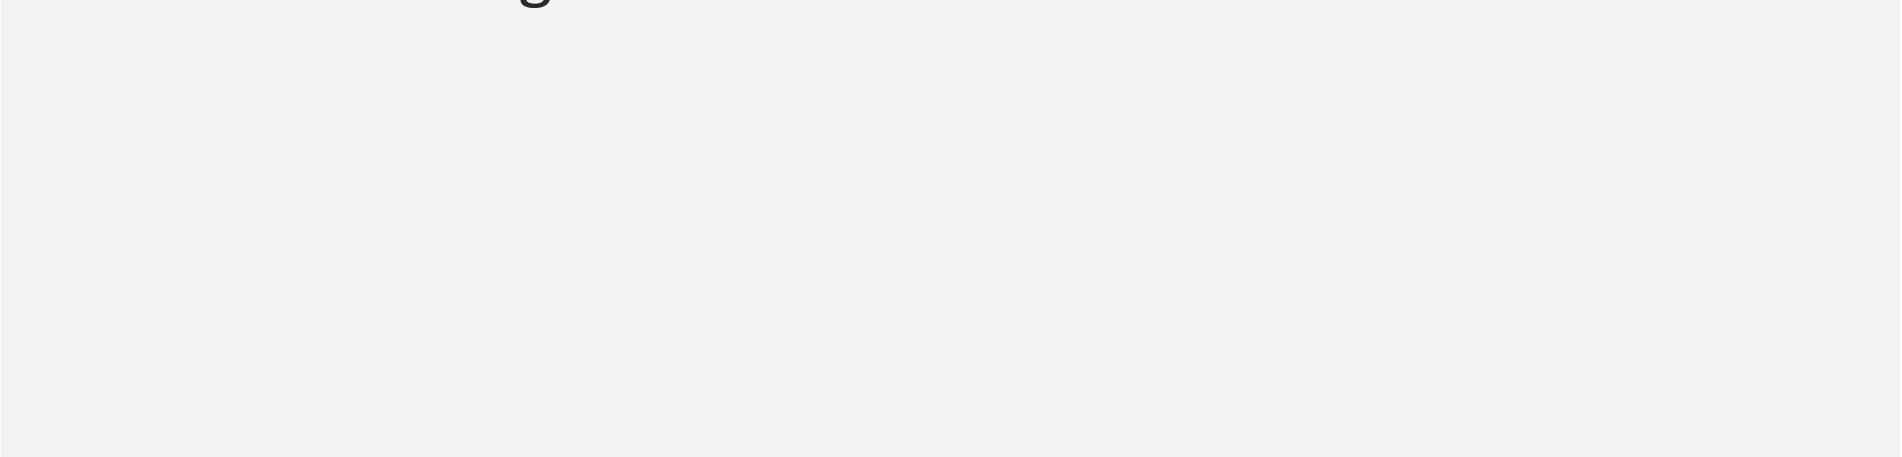
“Remember that all models are wrong; the practical question is how wrong do they have to be to not be useful.” George Box, 1987.

Assessing the loss




Assessing the loss


Part 1: Training error




Define training data



Define training data



Example:
Fit quadratic to minimize RSS




Training error (\hat{w}) =

$$\frac{1}{N} \sum_{i=1}^N (y_i - f_{\hat{w}}(x_i))^2$$

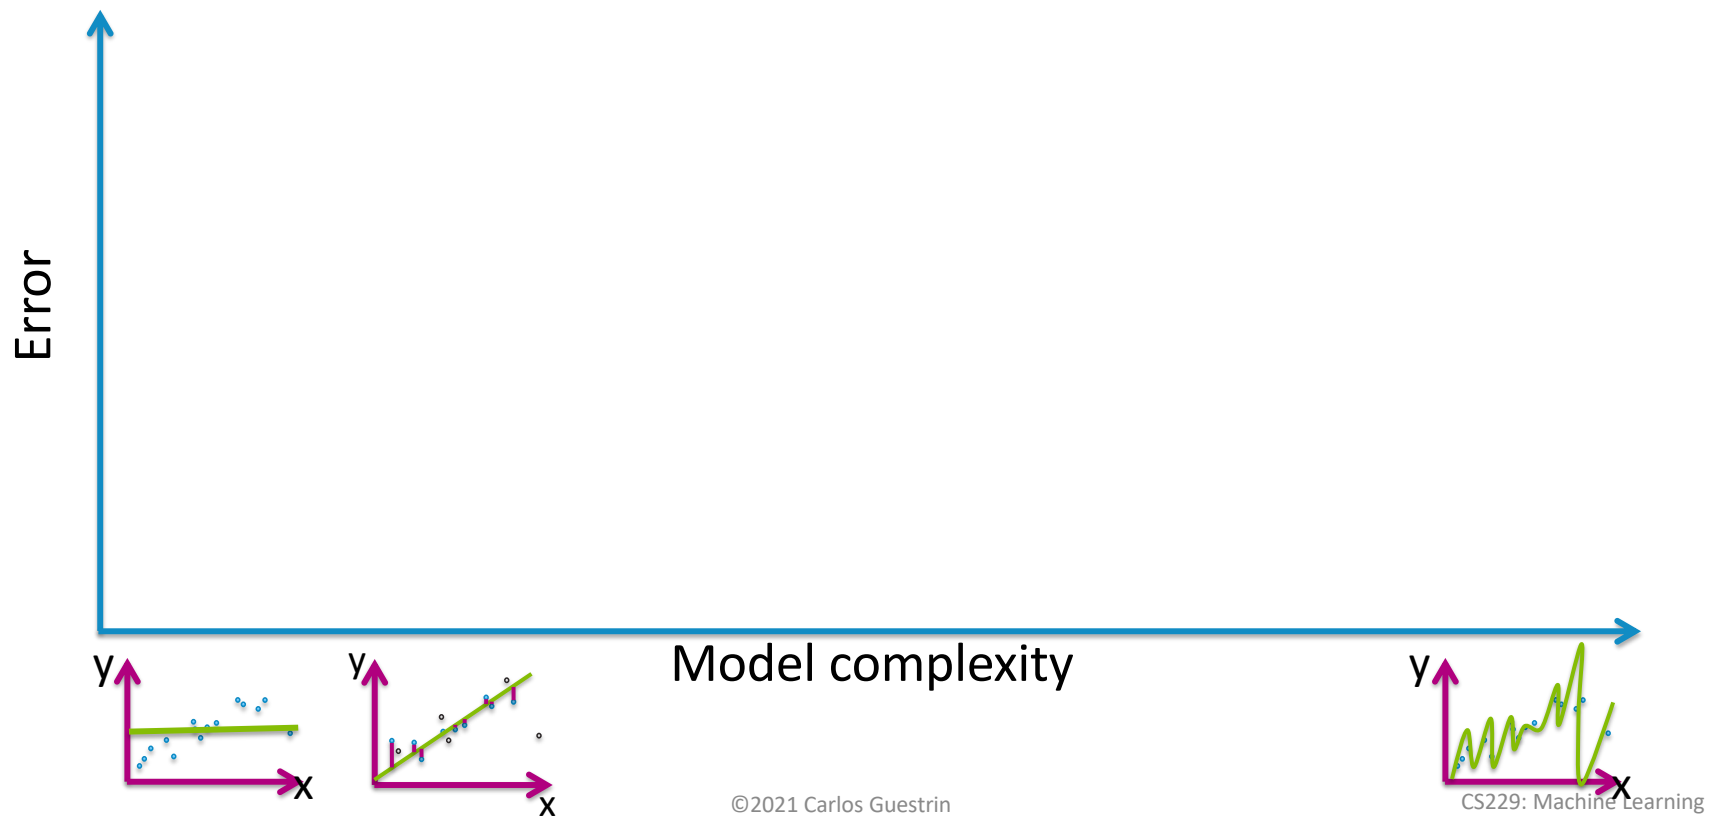
Example:

Use **squared error** loss $(y - f_{\hat{w}}(x))^2$



Training error (\hat{w}) = $\frac{1}{N} * [(\$_{\text{train } 1} - f_{\hat{w}}(\text{sq.ft.}_{\text{train } 1}))^2 + (\$_{\text{train } 2} - f_{\hat{w}}(\text{sq.ft.}_{\text{train } 2}))^2 + (\$_{\text{train } 3} - f_{\hat{w}}(\text{sq.ft.}_{\text{train } 3}))^2 + \dots \text{include all training houses}]$


Training error vs. model complexity




Is training error a good measure of predictive performance?

Issue:

Training error is overly optimistic... \hat{w} was fit to training data

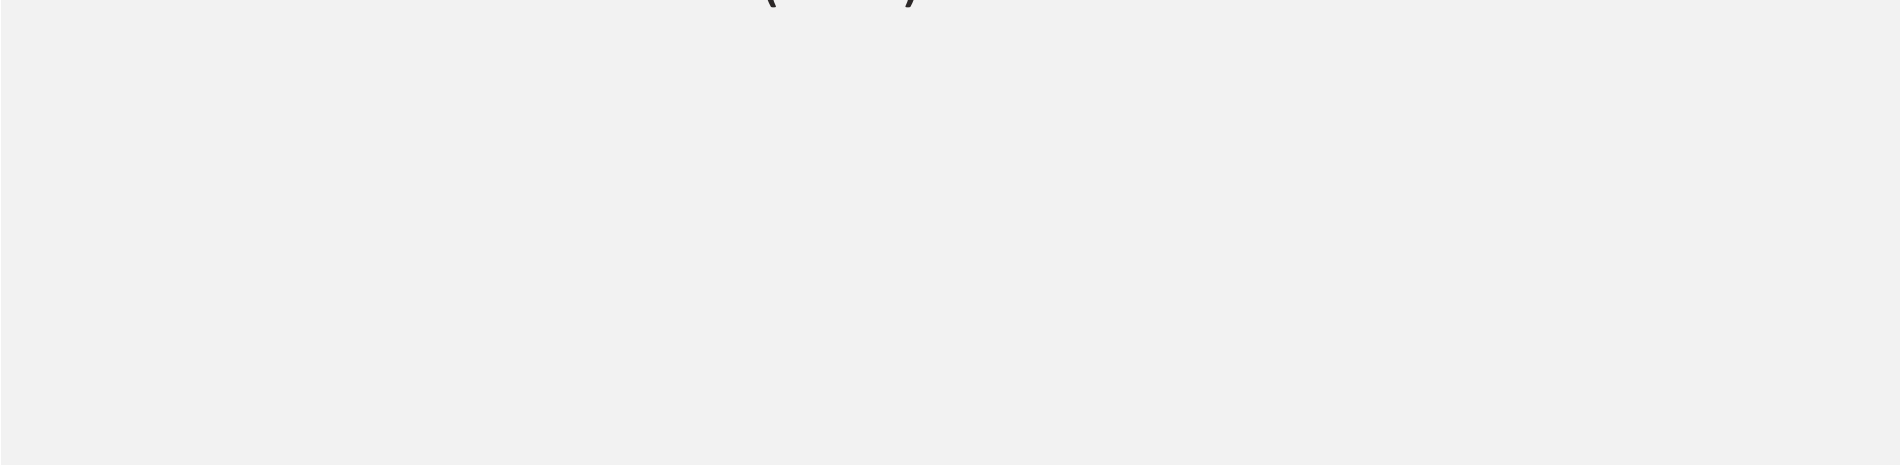


Small training error \neq good predictions
(unless training data includes
everything you might ever see)



Assessing the loss

Part 2: Generalization (true) error



Generalization error

Really want estimate of loss over all possible (, ) pairs



Generalization error definition

Really want estimate of loss over all possible (🏠,\$) pairs




Formally:

average over all possible
(x,y) pairs weighted by
how likely each is


$$\text{generalization error} = E_{x,y} [L(y, f_{\hat{w}}(x))]$$


fit using training data

Generalization error vs. model complexity



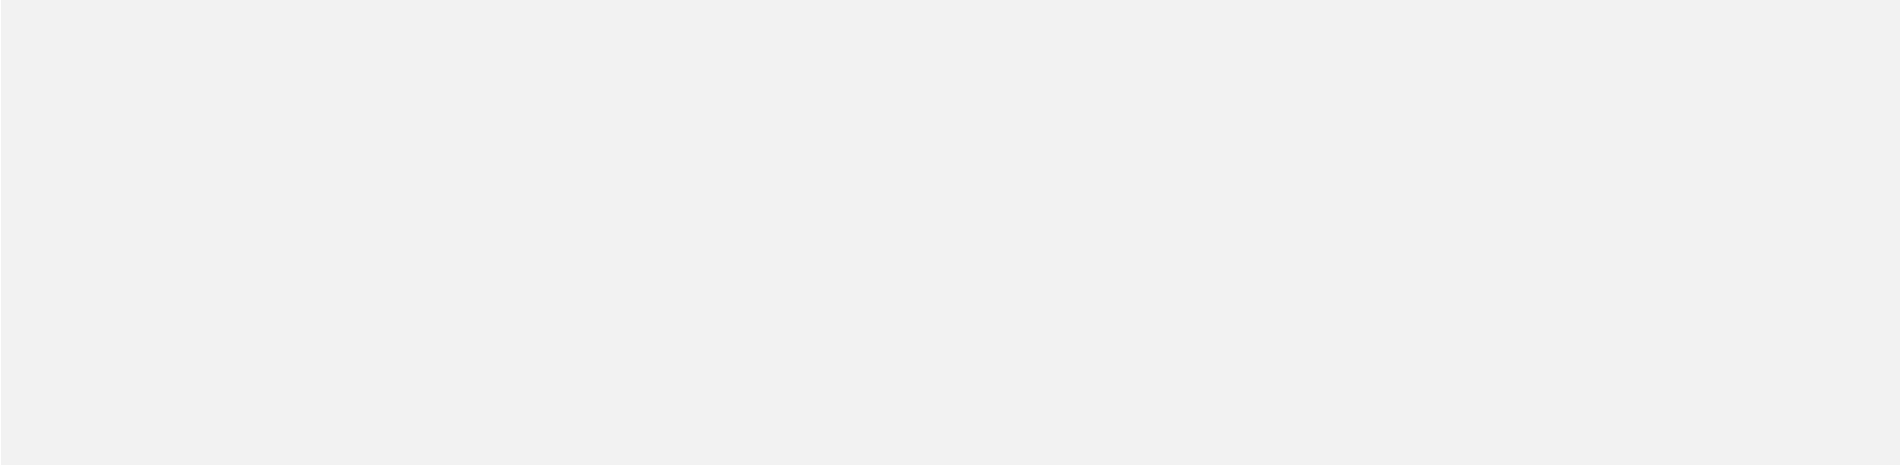
True error vs. model complexity






Assessing the loss

Part 3: Test error



Training, true, test error vs. model complexity

Overfitting if:




3 sources of error + the bias-variance tradeoff

3 sources of error

In forming predictions, there are 3 sources of error:

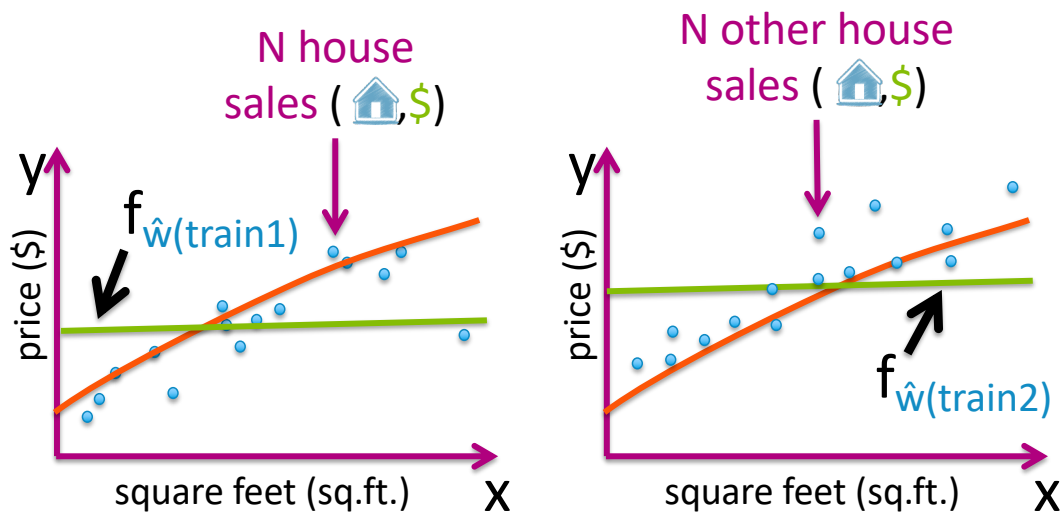
1. Noise
2. Bias
3. Variance

Data inherently noisy




Bias contribution

Suppose we fit a constant function



Bias contribution


Over all possible size N training sets, what do I expect my fit to be?



Bias contribution

$$\text{Bias}(x) = f_{w(\text{true})}(x) - f_{\bar{w}}(x)$$


← Is our approach flexible enough to capture $f_{w(\text{true})}$?
If not, error in predictions.



low complexity
→


Variance contribution

How much do specific fits vary from the expected fit?




Variance contribution

How much do specific fits vary from the expected fit?




Variance contribution

How much do specific fits vary from the expected fit?




Variance of high-complexity models

Assume we fit a high-order polynomial




Variance of high-complexity models

Suppose we fit a high-order polynomial




Variance of high-complexity models



high complexity



Bias of high-complexity models




high complexity




Sum of 3 sources of error

Average squared error at x_t

$$= \sigma^2 + [\text{bias}(f_{\hat{w}}(x_t))]^2 + \text{var}(f_{\hat{w}}(x_t))$$



Bias-variance tradeoff



Error vs. amount of data



Summary of bias-variance tradeoff

What you can do now...

- Contrast relationship between model complexity and train, true and test loss
- Compute training and test error given a loss function for different model complexities
- List and interpret the 3 sources of avg. prediction error
 - Irreducible error, bias, and variance